



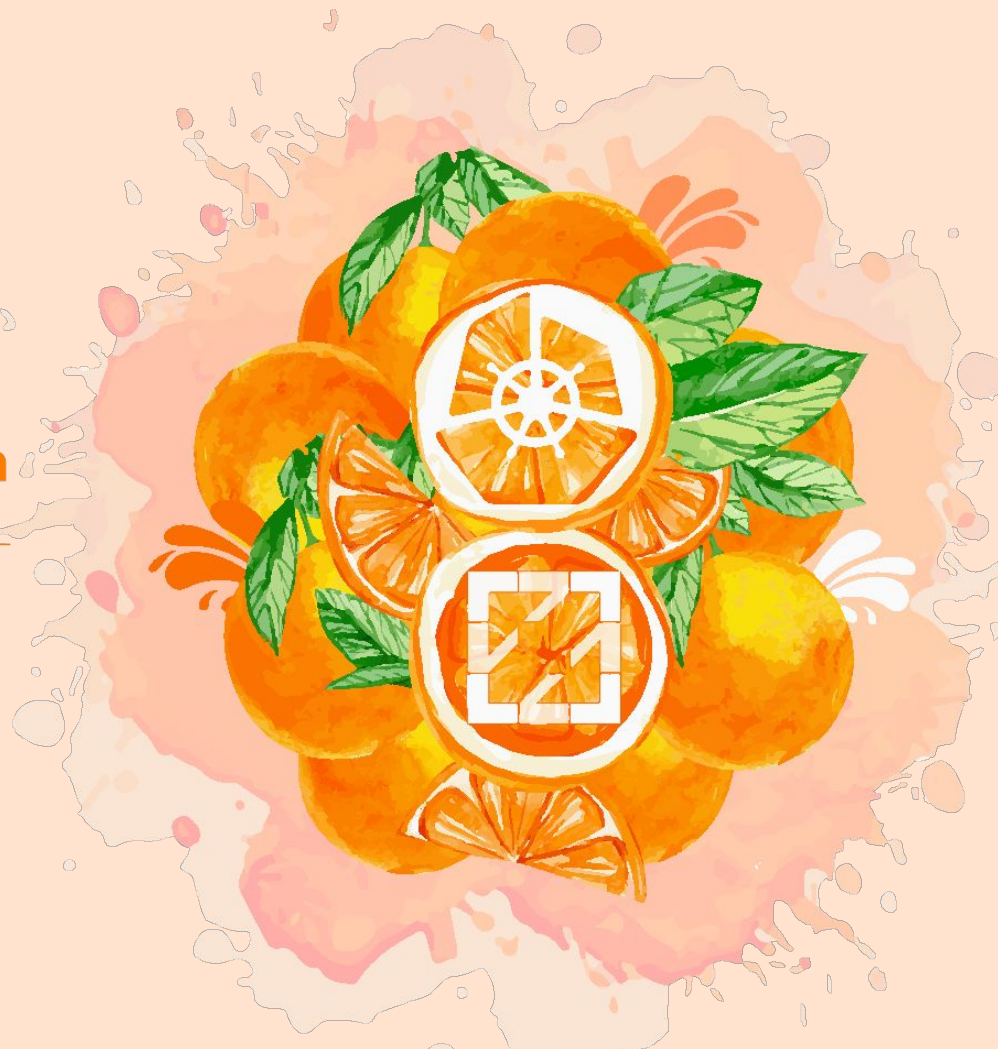
KubeCon



CloudNativeCon

Europe 2022

WELCOME TO VALENCIA





KubeCon



CloudNativeCon

Europe 2022

Autoscaling Kubernetes Deployments: A (Mostly) Practical Guide

Natalie Serrino, New Relic



About

 Software engineer / TLM at New Relic

 Contributor to Pixie

 Observability and performance

 Formerly worked in the data space



PromCon



Natalie Serrino

Content

🤔 What is Kubernetes autoscaling (and why is it useful)?

🕒 What knobs does Kubernetes autoscaling give us?

📈 Selecting an autoscaling metric for your application

🧐 A Turing-complete autoscaler (?!)

What is Kubernetes autoscaling?

Resource sizing in Kubernetes

How do you select the following values?



of nodes in your cluster



of pods in a deployment




Amount of resources to give a pod

Methods...

- Random guess
- Copypasta
- Proactive iteration
- Reactive iteration

Why is autoscaling useful?

- Ideal resource allocation depends on workload
- Workloads are often spikey and unpredictable



- Too few resources → bad user experience, latency, outages 🤖
- Too many resources → wasteful, expensive 💰 💰 💰

The types of autoscaling in Kubernetes

Kubernetes autoscaler	...adds/removes...	...to your...	...based on...
Cluster Autoscaler	Nodes	Cluster	Resource utilization
VerticalPodAutoscaler	Resources like CPU/Memory	Existing replicas	Resource utilization
HorizontalPodAutoscaler	Replicas	Workload	Resource utilization <u>or</u> user-defined metrics

Cluster Autoscaler



Set pod resource requests and limits



Make sure resource requests reflect actual usage



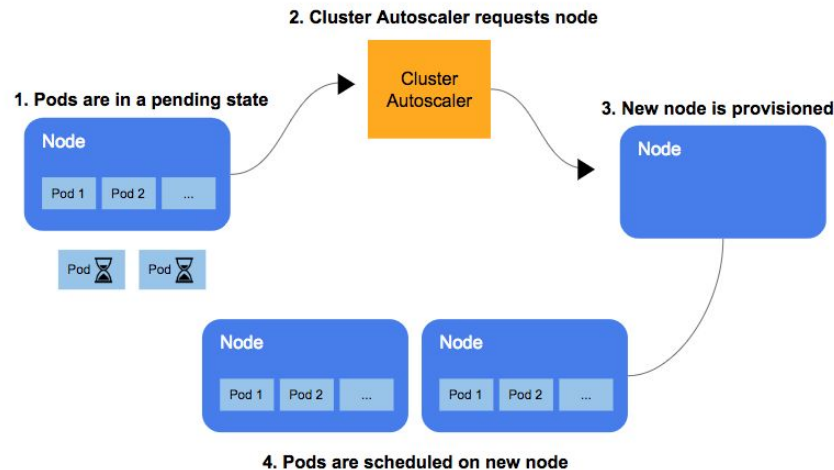
Specify PodDisruptionBudgets



Compare limits with available quota from your cloud provider (if applicable)



Kubernetes contributors tested for ≤ 1000 nodes with 30 pods each



*Image Credit: Understanding Kubernetes
Cluster Autoscaling by Ajay Tripathy*

VerticalPodAutoscaler (VPA)

- 💡 Can still set resource caps with VPA
- 💡 Updates may result in container restarts or pod rescheduling
- 💡 Use in cluster autoscaler to avoid VPA recommending more than available resources
- ⚠️ Can't use with HorizontalPodAutoscaler on the same application for CPU/memory yet
- ⚠️ Has not yet been tested on large clusters




Image Credit: Red Hat, How Full is my Cluster, Raffaele Spazzoli

HorizontalPodAutoscaler (HPA)



Lots of flexibility for metric selection




Check your service client affinity policies to ensure even load distribution



When scaling on CPU/memory, make sure to set resource requests



Can't use with VerticalPodAutoscaler on same application on CPU/memory yet



Demo: Horizontal pod autoscaling on CPU


HPA autoscaling equation

```
outputReplicas = ceil(  
    currentReplicas * ( currentMetricValue / desiredMetricValue )  
)
```



What knobs does Kubernetes autoscaling give us?

Scaling up and down



- How to set the minimum and maximum number of replicas?
- How often to look for changes in the metric?
- How quickly to add or remove pods?
- How many pods to add/remove in a single period?



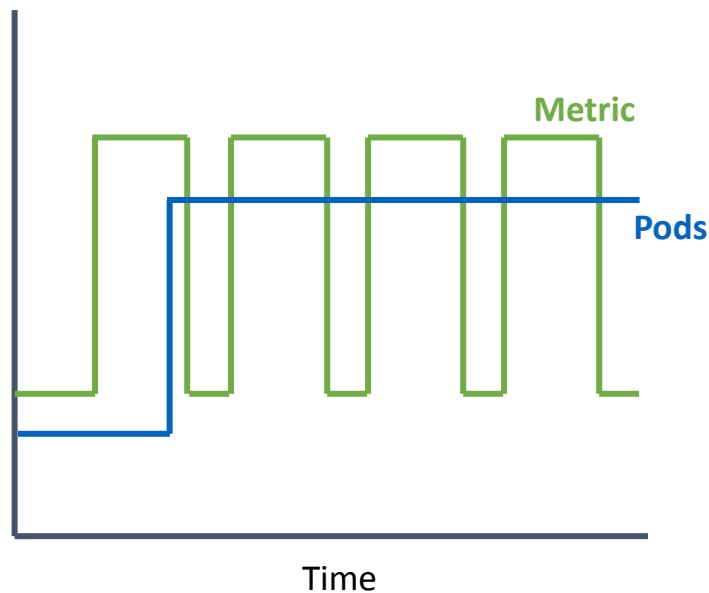

Capping min/max pods





Stabilization period




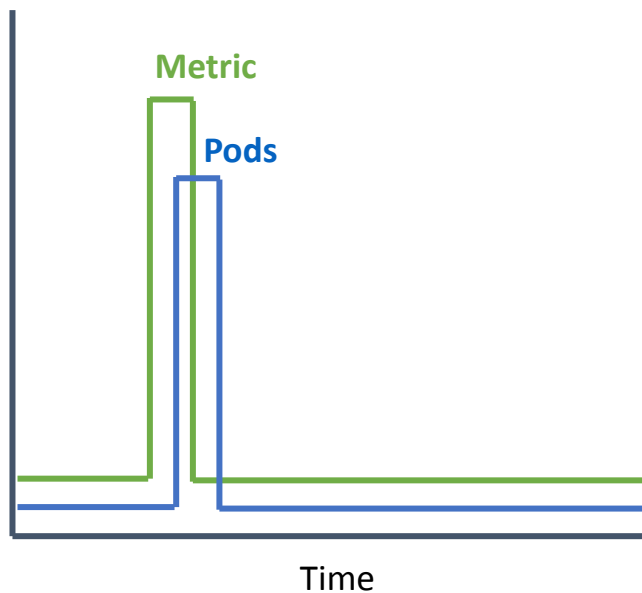
Stabilization period reduces pod churn



Capping max pods to add/remove per period



Capping max step size reduces pod churn



Selecting an autoscaling metric for your application

HPA Metric Types

Category	Type	Description	Examples
Resource Metrics	Built-in	Resource utilization metrics for pods and nodes only	Currently limited to CPU and memory
Custom Metrics	User-defined	Custom metrics about Kubernetes resources	Latency, throughput, queue depth
External Metrics	User-defined	Custom metrics NOT about Kubernetes resources	# of customers using website

Possible bottlenecks in an application


- CPU
- Memory
- Network
- # of worker threads
- # of outbound connections
- Downstream dependencies
- Queue depth
- ...Many more

The best metric to scale on depends on your workload!


Before

#1


Requests Per Second (Across Service)



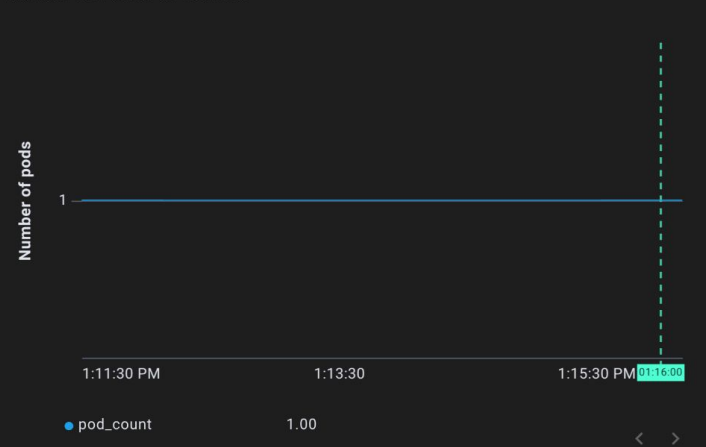
HTTP Latency



CPU Usage By Pod




Number Of Pods For Service




After #1


Requests Per Second (Across Service)




HTTP Latency



CPU Usage By Pod




Number Of Pods For Service




Before #2


Requests Per Second (Across Service)



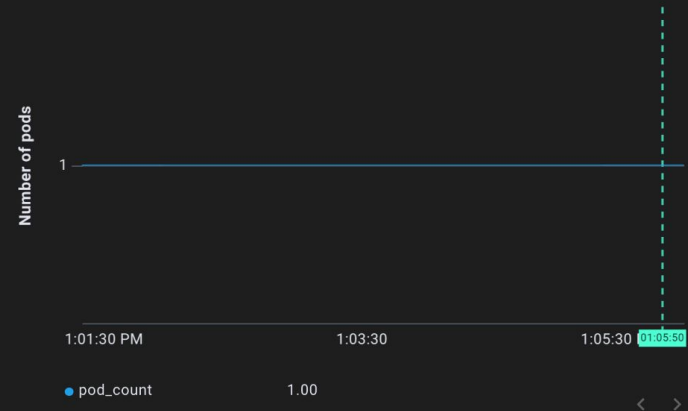
HTTP Latency



CPU Usage By Pod




Number Of Pods For Service




After #2


Requests Per Second (Across Service)




HTTP Latency



CPU Usage By Pod



Number Of Pods For Service



A Turing-complete autoscaler?

Turing machine

Turing machines are capable of any computation, given enough time and tape.

Input program

Instruction 1
Instruction 2
Instruction 3
...
Instruction N

*Next
instruction*

Turing Machine

Write value

Output tape

3

0

7

1

4

3

3

2

5

2



Subleq

“One-instruction set computer”, sufficient for Turing completeness.

Instruction **subleq** **a**, **b**, **c**

$\text{Mem}[\mathbf{b}] = \text{Mem}[\mathbf{b}] - \text{Mem}[\mathbf{a}]$

if ($\text{Mem}[\mathbf{b}] \leq 0$)

goto **c**

Subleq pseudocode (credit: Wikipedia)

HPA subleq Turing machine

Input program

```
subleq 3, 0, 3
subleq 5, 3, 2
subleq 2, 1, 1
...
subleq 4, -1, 0
```

*Execute 1
instruction
per HPA
interval*

Custom HPA
metric provider

Set # pods

Number of pods
over time

3

1

7

1

4

3

3

2

5

2



Input program: deployment name


9x-1x3x10x-1x6x0x0x-1x72x105x0

9, -1, 3, 10, -1, 6, 0, 0, -1, 72, 105, 0

```
subleq 9, -1, 3
subleq 10, -1, 6
subleq 0, 0, -1
subleq 72, 105, 0
```




Split on "x"



*3 input args
per subleq*

Setting a certain # of output pods

```
outputReplicas = ceil(  
    currentReplicas * ( currentMetricValue / targetMetricValue )  
)
```



*Backwards calulating
“current metric value”*

```
currentMetricValue = (  
    targetMetricValue * ( outputReplicas / currentReplicas )  
)
```

Demo: Turing-complete autoscaler



kubernetes

SIG instrumentation

github.com/kubernetes-sigs/custom-metrics-apiserver



Load generator

github.com/rakyll/hey



KubeCon



CloudNativeCon

Europe 2022



CNCF Sandbox Project

github.com/pixie-io/pixie
blog.px.dev

Thanks!

