



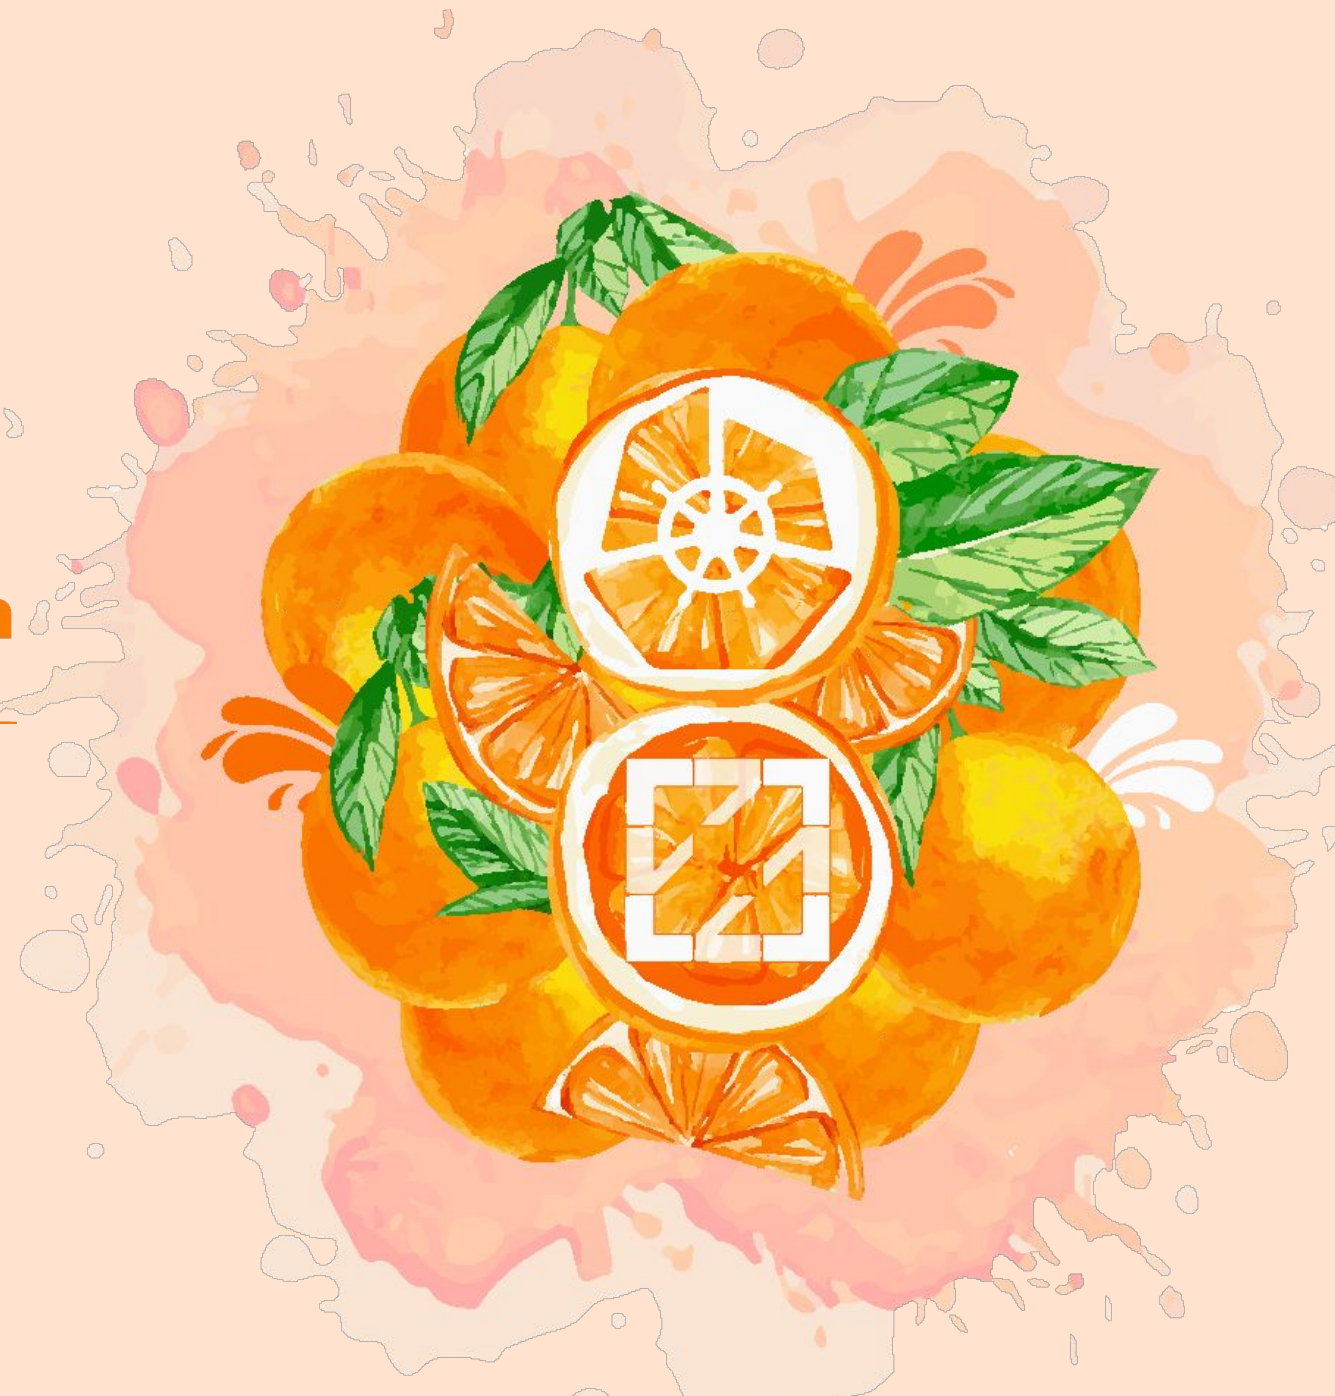
**KubeCon**



**CloudNativeCon**

**Europe 2022**

**WELCOME TO VALENCIA**





KubeCon



CloudNativeCon

Europe 2022

# Build your own Cluster API Provider

The ~~hard~~ easy way



# Who are we?



Anusha Hegde

Senior Engineer @ VMWare  
Cluster API Provider BYOH Maintainer



Richard Case

Principal Engineer @ Weaveworks  
Cluster API Provider AWS Maintainer  
Cluster API Provider Microvm Maintainer

Lego & retro gaming addict

Who uses Cluster API (CAPI) already?

Who contributes to a Cluster API Provider?

Who's thinking of building a Cluster API Provider?

# What will we be covering?

- What is CAPI
- Provider Types
- Design
  - First rule of creating a provider.....you don't create a provider. Do you really need one???
  - API design, versioning
  - Controller and CAPI contracts
  - Webhooks
- Patterns
  - Experimental / feature flags
- Dev
  - Tools - kubebuilder, tilt, ...
  - Project layout - conventions used by capi & kubebuilder
  - Debugging...tilt/delve is your friend (vscode, goland with tilt)
- Test
  - E2e -CAPI e2e test framework, tradition testing pyramid is out the window
  - Prow when you donate (more later on dnation)
- Community
  - Donation, office hours, slack etc

# What will we be covering?

- What is Cluster API
- Different Provider types
- Designing a Provider
- Common Patterns
- Development & Testing
- Community



KubeCon



CloudNativeCon

Europe 2022

# What is Cluster API? (1/2)

- Built on the premise that “Cluster lifecycle management is difficult”
- Declarative specification of clusters
- Establishes building blocks for higher order functionality
  - Cluster templating
  - Automation of scaling, repair & upgrades
  - Distributing nodes across failure domains
  - Machine health checks to replace unhealthy nodes
  - Managed control planes
  - .....

# What is Cluster API? (2/2)

- Designed around interchangeable components via “providers”
- **clusterctl** handles the lifecycle of a CAPI management cluster
  - Cluster templates / flavors are very useful
  - Providers operator being developed.....GitOps friendly provider operations :)
- Community calls every week on Wednesday @ 6pm GMT / 10am PT
  - Separate calls for providers
- For a walkthrough of CAPI see the “lets talk about...” series by Stefan:
  - <https://github.com/kubernetes-sigs/cluster-api/discussions/6106>



# What is Cluster API?

- Built on the premise that “Cluster lifecycle management is difficult”
- Declarative specification of clusters
- **clusterctl** handles the lifecycle of a CAPI management cluster
- Community calls every week on Wednesday @ 6pm GMT / 10am PT
- For a walkthrough of CAPI see the “lets talk about...” series by Stefan:
  - <https://github.com/kubernetes-sigs/cluster-api/discussions/6106>

# What is a Cluster API provider?

A Kubernetes **operator** that implements infrastructure / operating environment specific functionality that is utilized by core Cluster API when managing the lifecycle of a K8s cluster.

The operator implements a contract via its custom resources (i.e. CRDs) depending on the type of provider, which enables interaction between core CAPI and the provider.


# Provider Types

- **Infrastructure** - used to provision any infrastructure that is required to create and run a Kubernetes cluster. For example, networking, security groups, virtual or physical host machines
- **Bootstrap** - used to create the “user-data” that is passed to the infrastructure machines that contains the instructions to bootstrap a Kubernetes node on that machine. 2 parts to it:
  - Action: how Kubernetes is bootstrapped (e.g. invoking kubeadm)
  - Format: how the action is encoded and passed to the machine (e.g. cloud-init, ignition)
- **Control plane** - used to control the creation & lifecycle of the Kubernetes control plane. It can utilize resources created by bootstrap and infrastructure providers.
  - Kubeadm control plane is the original
  - Managed Kubernetes (i.e. EKS, AKS) implementations - no nodes

# First rule of creating a provider...

## ...you don't need to create a provider!

(hopefully)



# What constitutes a Cluster API provider?

- A Kubernetes operator (a.k.a controller manager)
  - CRDs
  - Controllers that reconcile the CRDs
- k8s resources to deploy the controller
  - Plain old yaml
  - (Optional) tokens that will be replaced an installation time
  - Kustomize
- Metadata / repo layout



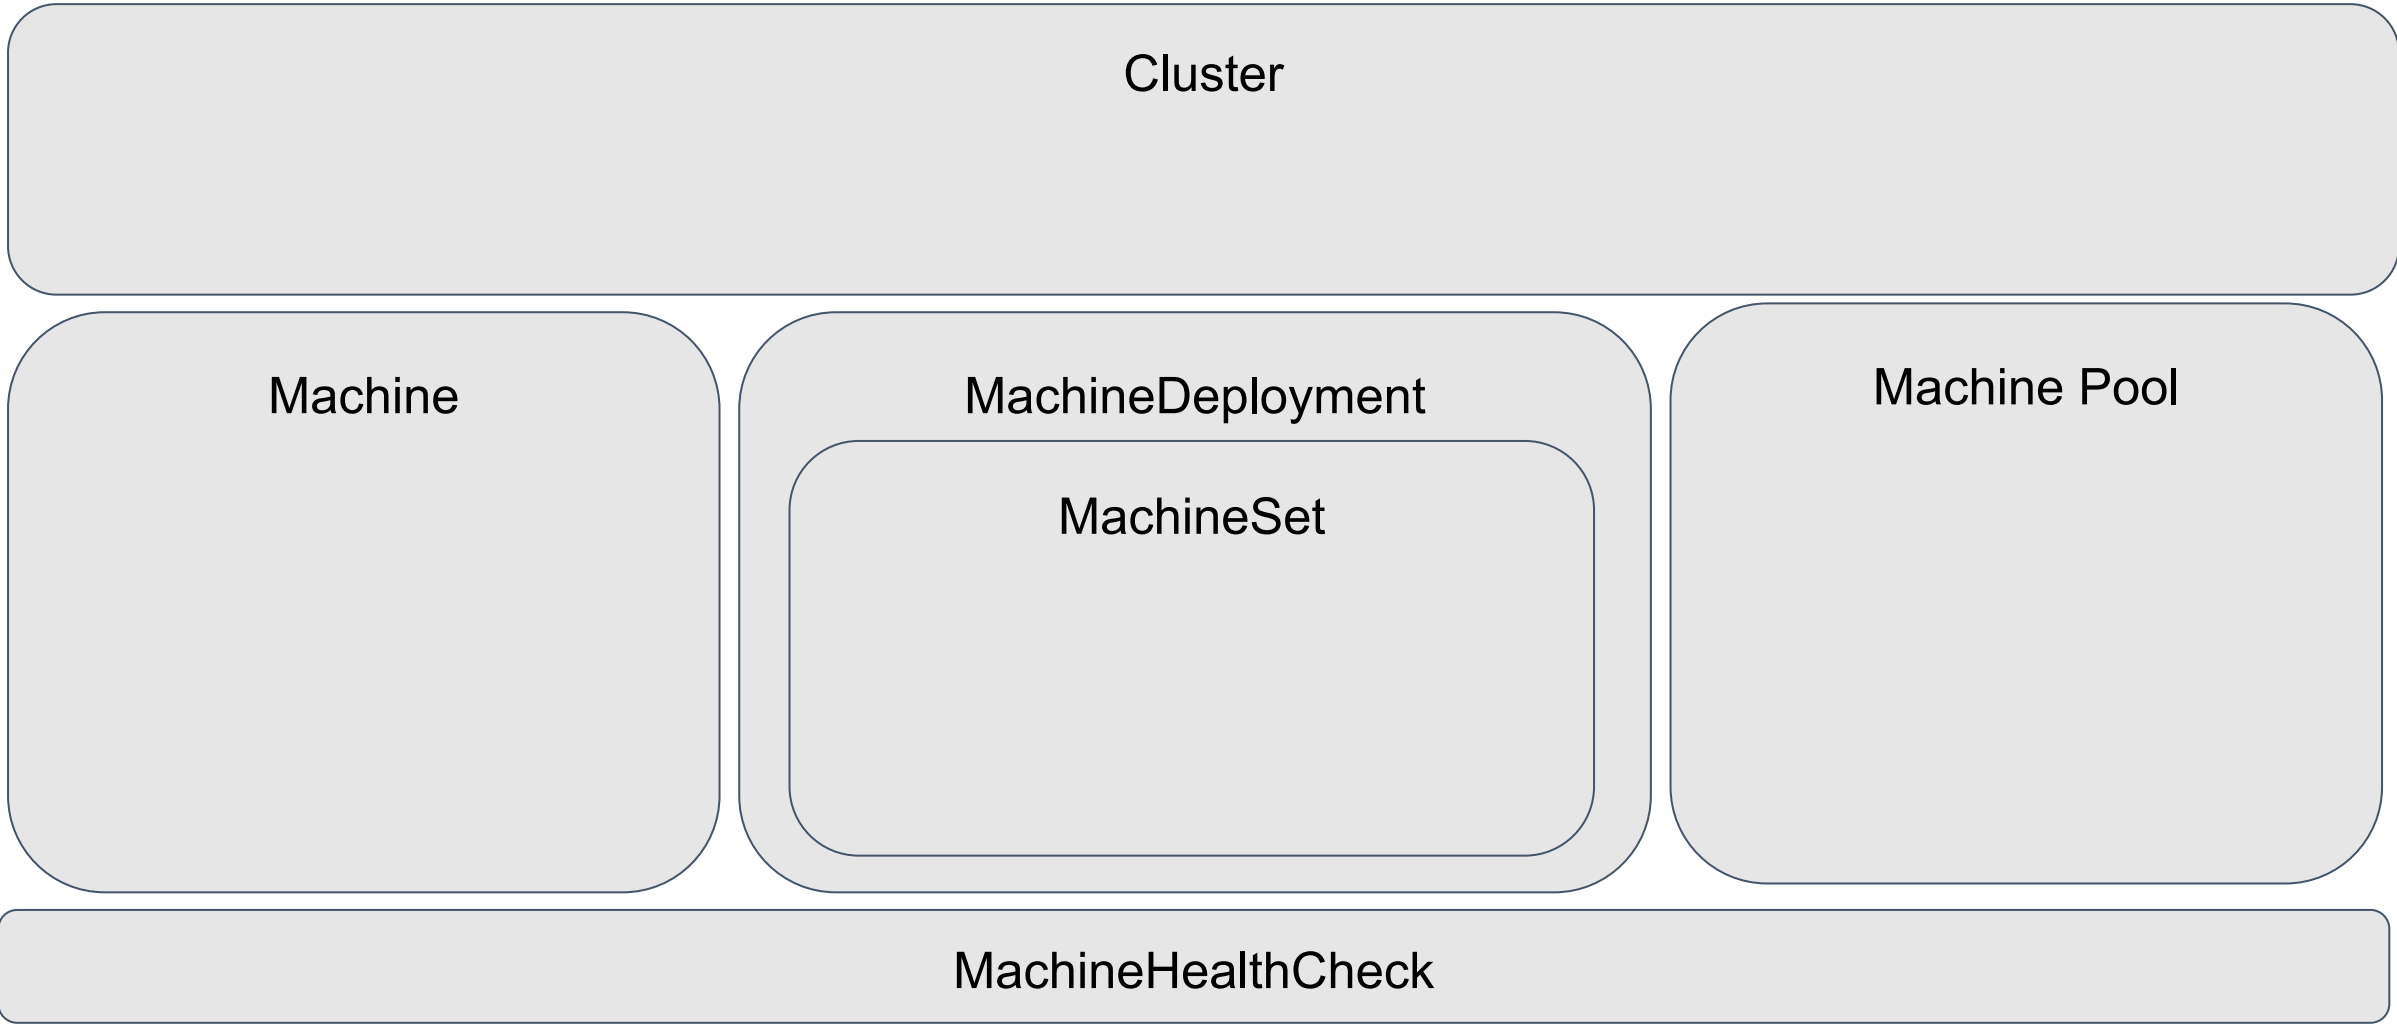
KubeCon



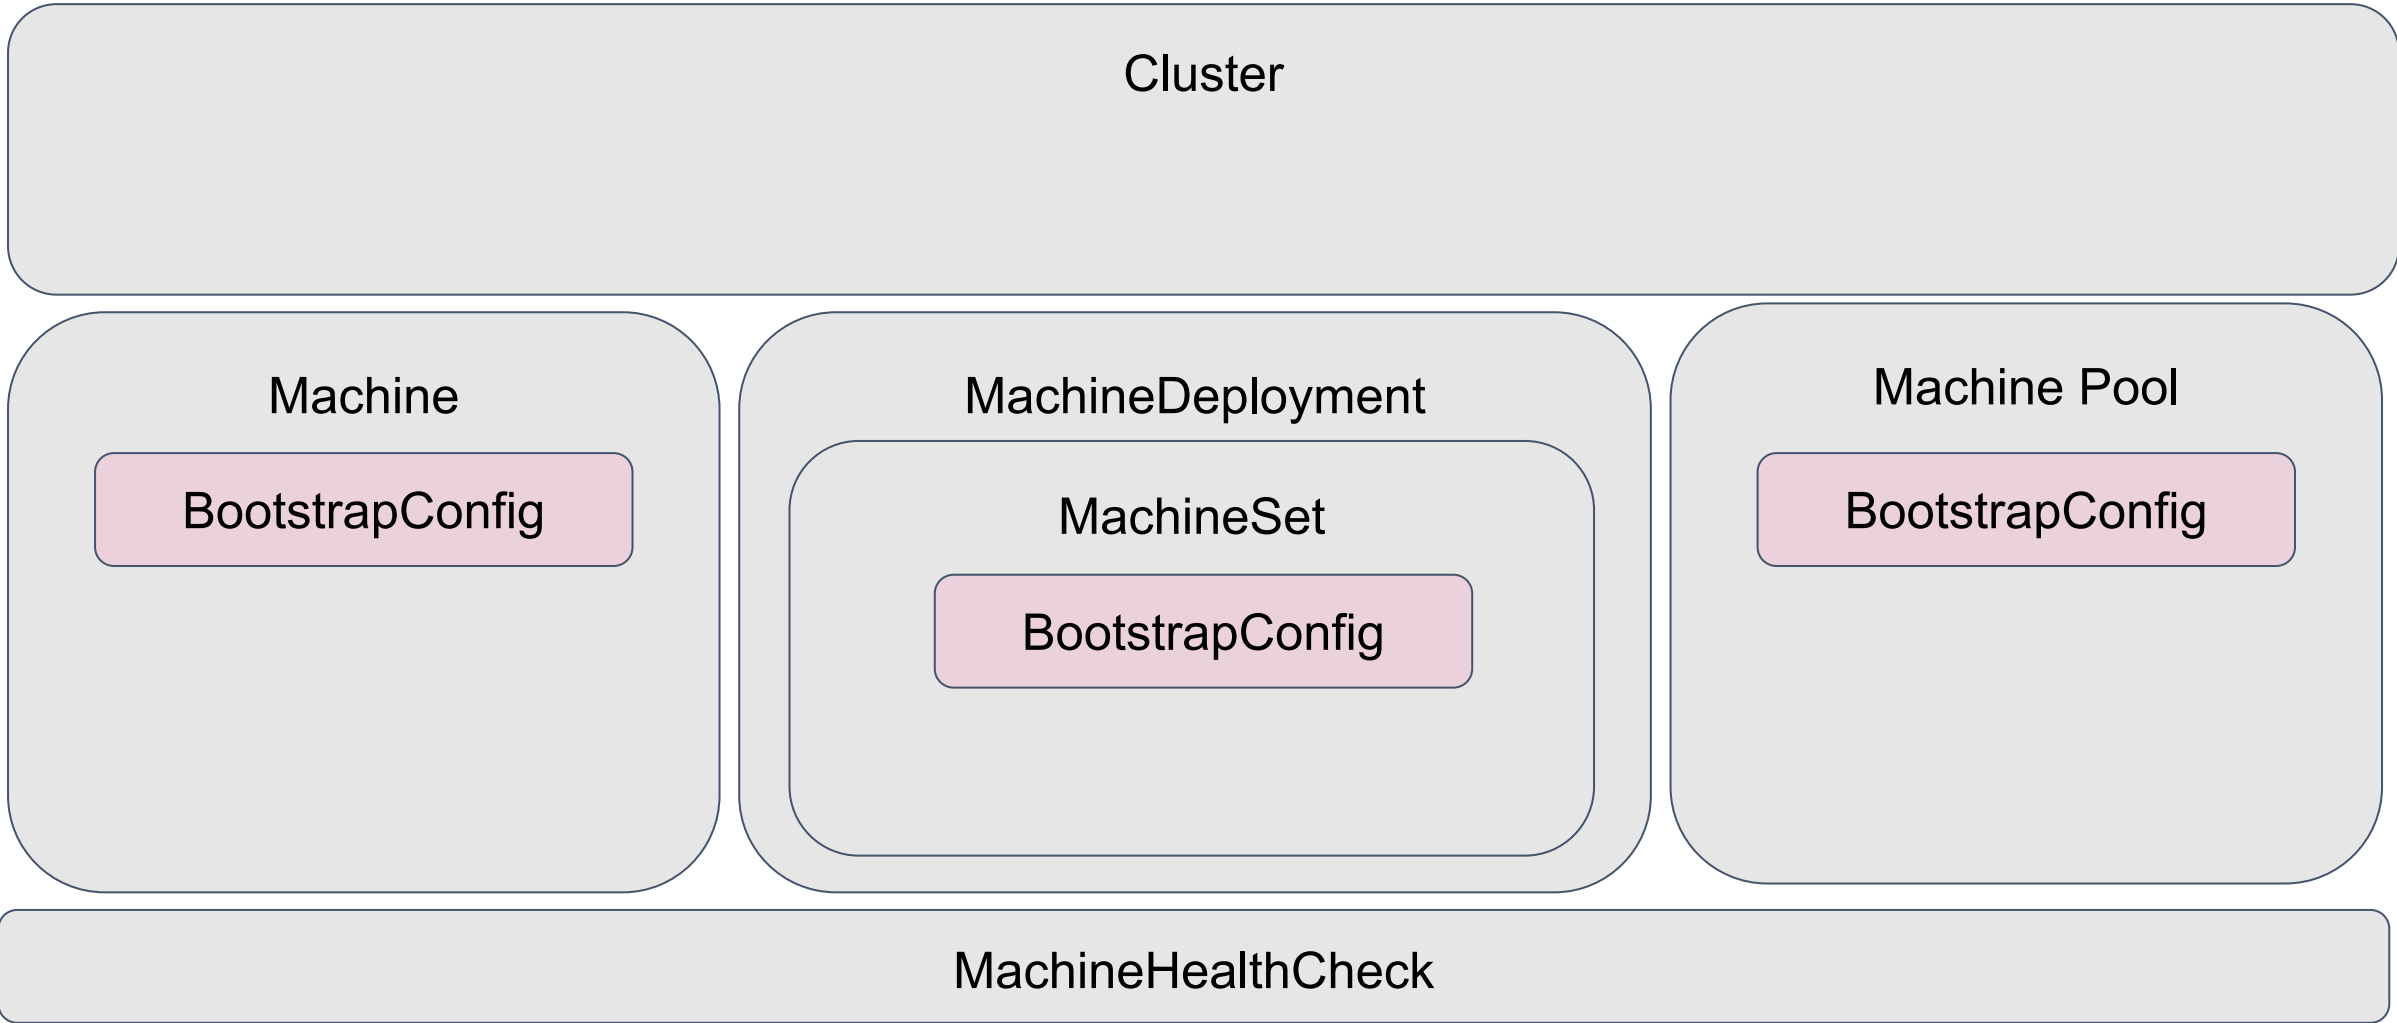
CloudNativeCon

Europe 2022

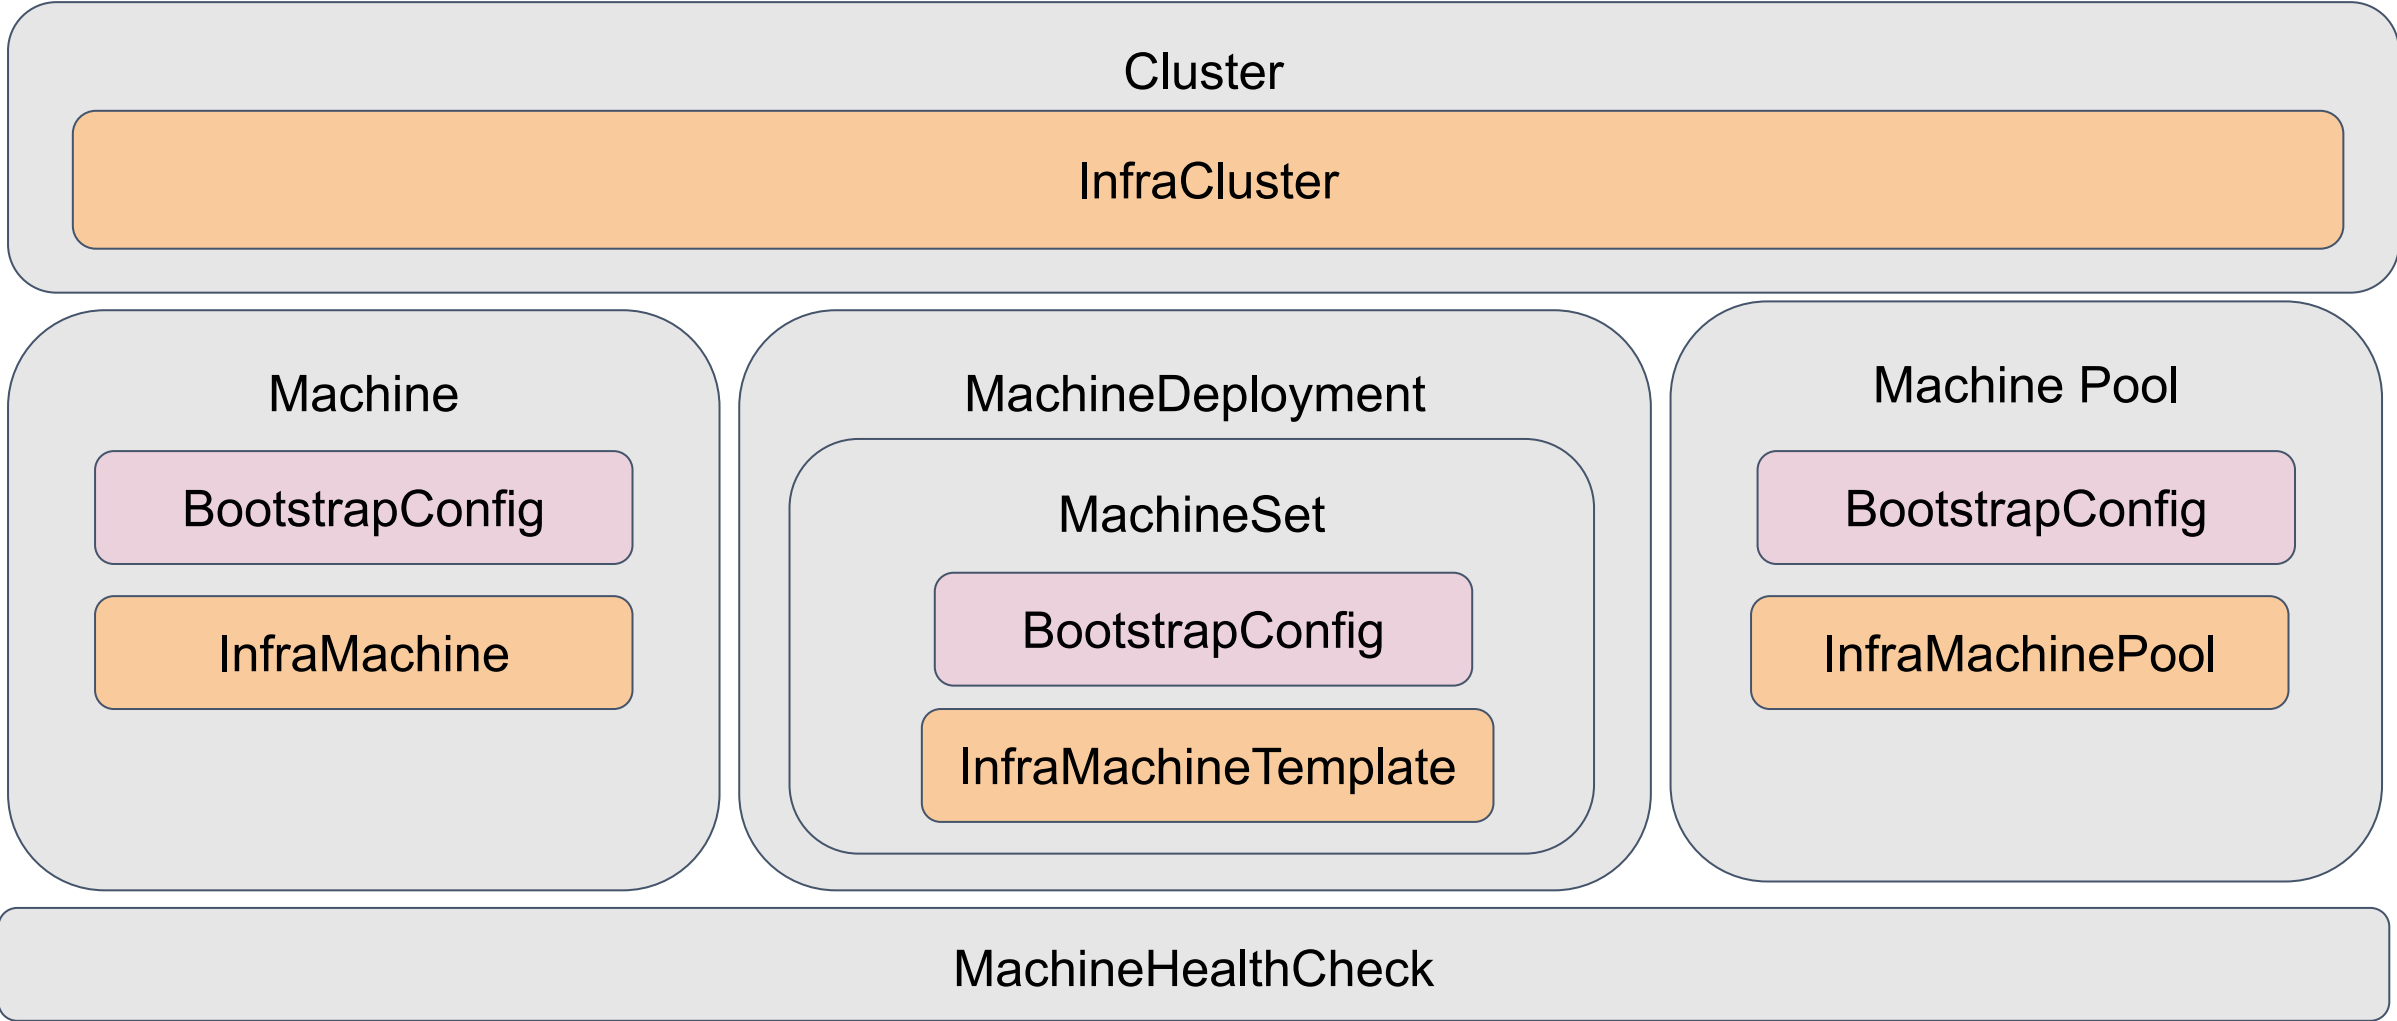
# Core



# Bootstrap

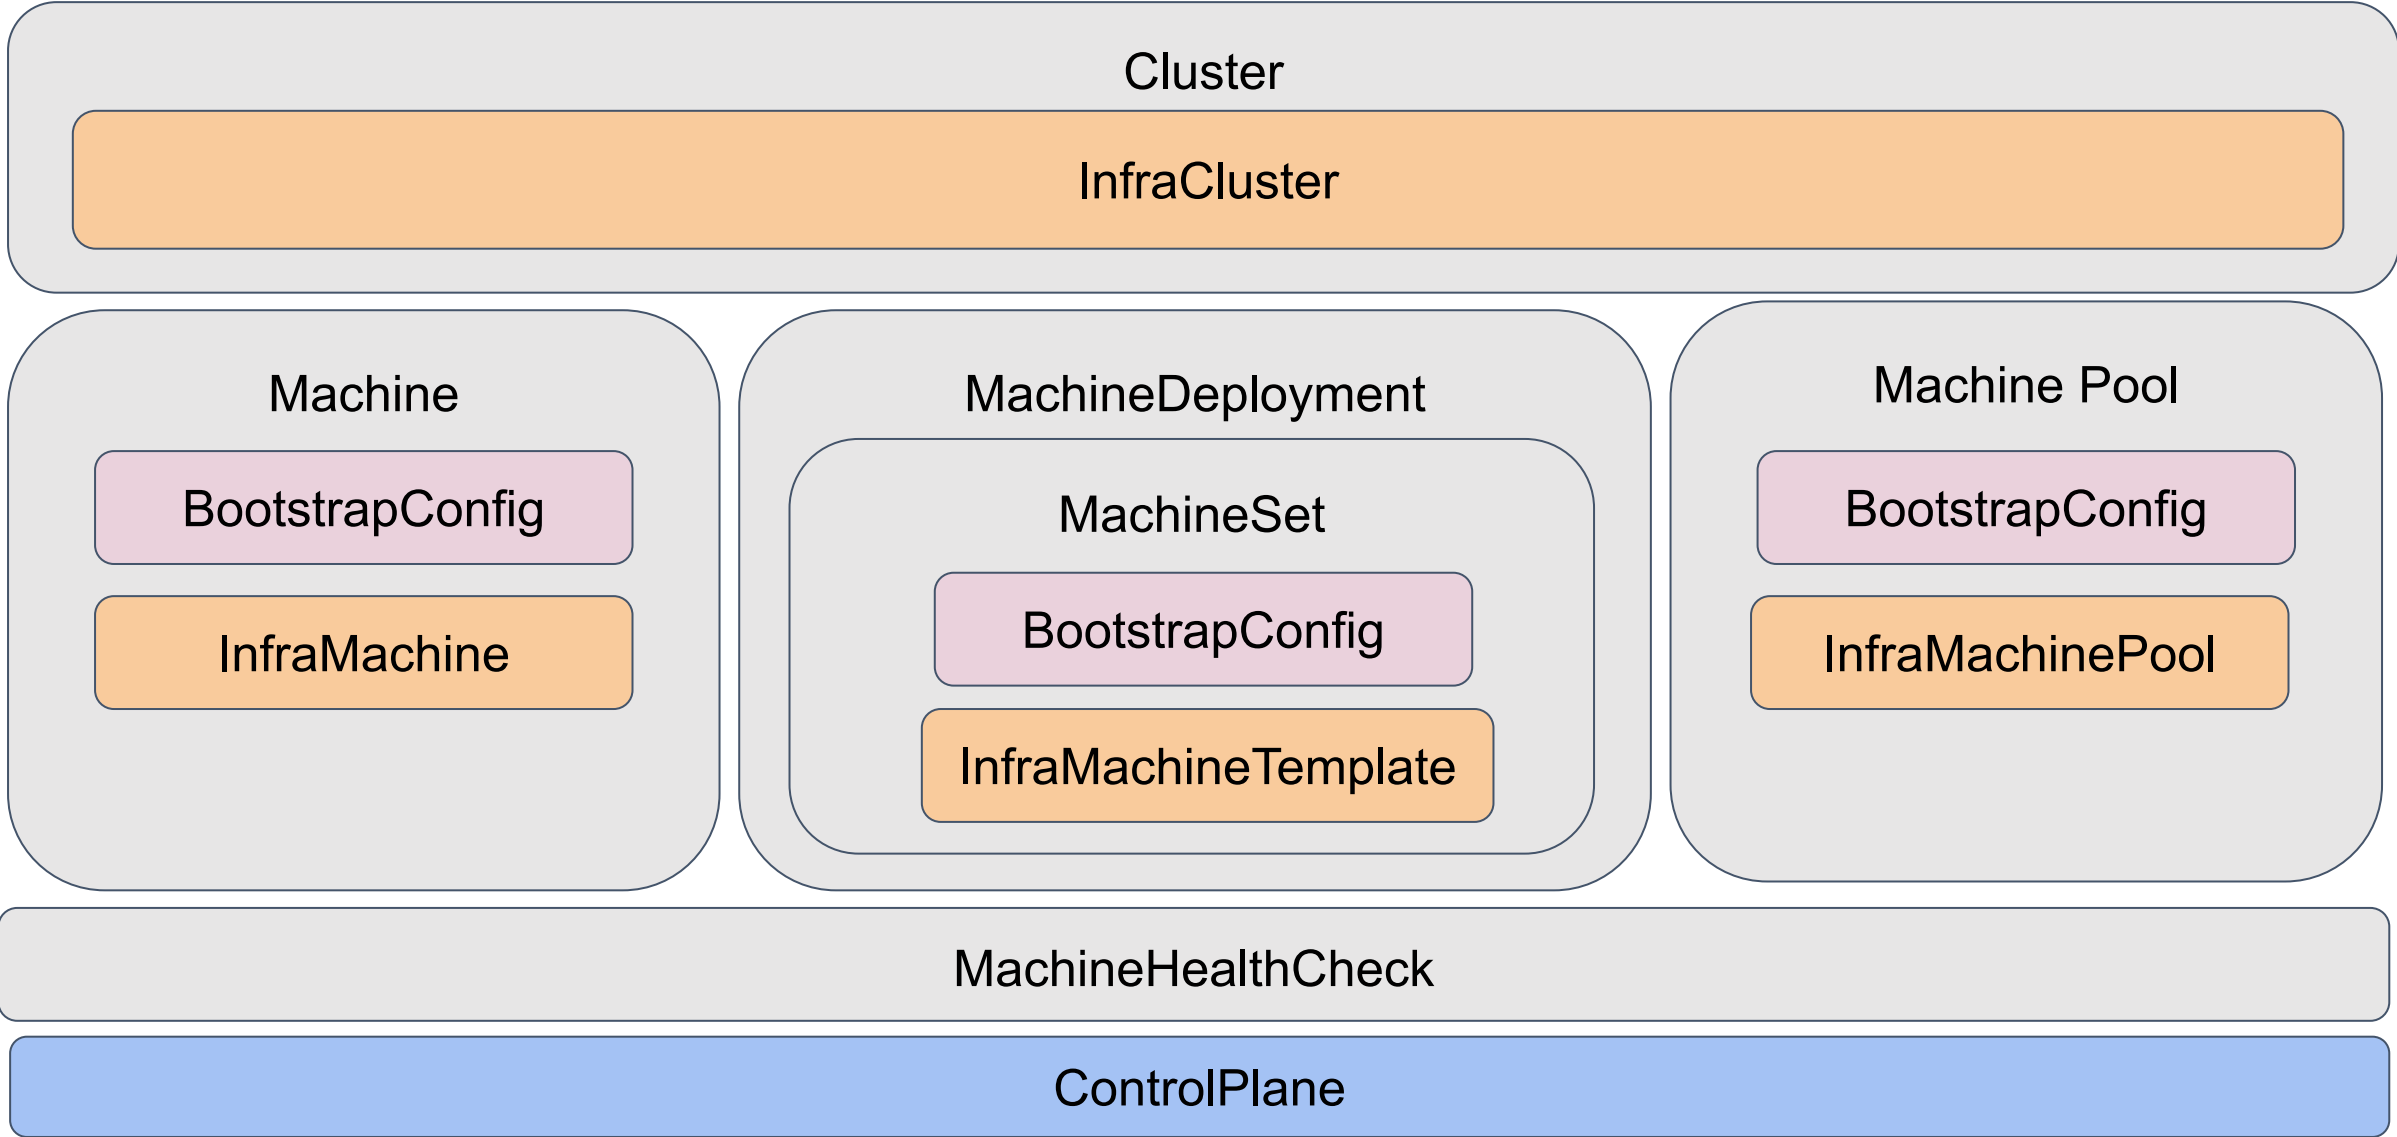


# Infrastructure





# Control Plane




# What is an operator?

- A way to create, manage and configure complex applications in Kubernetes.
- Codifies the steps a human would do to deploy and operate a complex application
  - For example, the steps to create Kubernetes cluster involved creating infrastructure, bootstrapping k8s, managing version upgrades.
- Surface to user via declarative API (i.e. CRDs)
- Contains one or more controllers that understand & reconcile the CRDs

# What is a controller?

“ A controller is a control loop that watches the desired state of the cluster through the API server and makes changes attempting to move the current state towards the desired state. “

# Control Loop



- Watch - for changes in custom resource(s)
- Diff - work out the difference between desired & actual state
- Act - take action to remediate the difference (if any)

.....And repeat!

# What provider type do you need?

- Do you operate a cloud / baremetal service?
  - You'll need an Infrastructure provider
- Do you want a different way to bootstrap Kubernetes instead of Kubeadm?
  - You'll need a bootstrap provider and maybe a control plane provider
- Do you have a hosted Kubernetes control plane service?
  - You may need a control plane provider
- Do you want to use virtualization technology like KVM or vSphere?
  - You may be covered by the vSphere, Microvm or Kubevirt providers
  - If not then may need to create a infrastructure provider
- Do you want to provision your own infrastructure and get CAPI to manage Kubernetes?
  - You may be covered by one of the existing providers that allow you to prong your own infrastructure (like CAPA, CAPZ)
  - You may be covered by the “bring your own host” provider
  - If not then you may need any of the 3 types of provider

# Scaffolding the provider (1/2)

- Kubebuilder and controller-runtime are your friends:

```
kubebuilder init --domain cluster.x-k8s.io --repo github.com/capi-samples/cluster-api-provider-podman
kubebuilder create api --group infrastructure --version v1alpha1 --kind PodmanCluster
kubebuilder create api --group infrastructure --version v1alpha1 --kind PodmanMachine
```

Important parts:

- **--domain cluster.x-k8s.io** - this generally used as the domain part of the GVK by providers
- **--group infrastructure** - this by convention indicates that the CRD relates to an infrastructure provider
- **--version v1alpha1** - providers need to follow the Kubernetes API versioning standard and the guarantees these provide to the user
- **--kind PodmanCluster/Machine** - by convention, a provider uses a consistent prefix on the CRD names

## Scaffolding the provider (2/2)

- We also need a machine template API type **BUT no controller**
- Add a reference to CAPI so we can use their API definitions / utility functions

```
→ kubebuilder create api --group infrastructure --version v1alpha1 --kind PodmanMachineTemplate
Create Resource [y/n]
y
Create Controller [y/n]
n
→ go get sigs.k8s.io/cluster-api@v1.1.2
```

Why do we not need a controller for the machine template?

- It's used as a template to create new instances of **PodmanMachine**
- **PodmanMachine** has a controller to handle reconciliation

# Provider Metadata

A provider must specify which versions are compatible with which CAPI API version.

- One of the requirements to be installable via **clusterctl init**
- Create a metadata.yaml file in the root of the repo

```
# maps release series of major.minor to cluster-api contract version
# the contract version may change between minor or major versions, but *not*
# between patch versions.
#
# update this file only when a new major or minor version is released
apiVersion: clusterctl.cluster.x-k8s.io/v1alpha3
releaseSeries:
- major: 0
  minor: 1
  contract: v1beta1
```



# Define API for your provider (1/2)



KubeCon



CloudNativeCon

Europe 2022

- Add fields to the **Spec & Status** of your API types to conform to your provider types contract
  - CAPI documentation helps: <https://cluster-api.sigs.k8s.io/developer/providers/implementers.html>

```
// PodmanClusterSpec defines the desired state of PodmanCluster
type PodmanClusterSpec struct {
    // ControlPlaneEndpoint represents the endpoint used to communicate with the control plane.
    //
    // See https://cluster-api.sigs.k8s.io/developer/architecture/controllers/cluster.html
    // for more details.
    //
    // +optional
    ControlPlaneEndpoint clusterv1.APIEndpoint `json:"controlPlaneEndpoint"`
}

// PodmanClusterStatus defines the observed state of PodmanCluster
type PodmanClusterStatus struct {
    // Ready indicates that the cluster is ready.
    // +optional
    // +kubebuilder:default=false
    Ready bool `json:"ready"`

    // FailureDomains is a list of the failure domains that CAPI should spread the machines across. For
    // the CAPPOD provider this doesn't mean anything.
    FailureDomains clusterv1.FailureDomains `json:"failureDomains,omitempty"`
}
```

## Define API for your provider (2/2)



KubeCon



CloudNativeCon

Europe 2022

- Add custom fields to **the Spec & Status** of your API types that are specific to your provider

```
// PodmanMachineSpec defines the desired state of PodmanMachine
type PodmanMachineSpec struct {
    // ProviderID will be the container name in ProviderID format (podman:///<containername>)
    // +optional
    ProviderID *string `json:"providerID,omitempty"`

    // ExtraMounts describes additional mount points for the node container
    // These may be used to bind a hostPath
    // +optional
    ExtraMounts []Mount `json:"extraMounts,omitempty"`
}
```

# Add finalizers

- **If your provider creates external resources you will need to define finalizers**
  - A finalizer allows the controller to clean up external resources before allowing the API type to be deleted from API server.
  - See docs: <https://book.kubebuilder.io/reference/using-finalizers.html>

```
const (  
    // MachineFinalizer allows ReconcilePodmanMachine to clean up resources associated with  
    // PodmanMachine before removing it from the apiserver.  
    MachineFinalizer = "podmanmachine.infrastructure.cluster.x-k8s.io"  
)
```

- The controllers will add and remove the finalizers

# Implement controllers for your API types (1/2)

- Kubebuilder will have created an “empty” controller with
  - **Reconcile** function
  - Controller setup to watch its CRD type
- Setup tasks that you need to do:
  - Watch companion CRDs from CAPI
  - Ensure reconciliation doesn't occur if paused or if the resource is externally managed

```
// SetupWithManager sets up the controller with the Manager.
func (r *PodmanMachineReconciler) SetupWithManager(ctx context.Context, mgr ctrl.Manager, options
controller.Options) error {
    log := ctrl.LoggerFrom(ctx)

    builder := ctrl.NewControllerManagedBy(mgr).
        WithOptions(options).
        For(&infrav1.PodmanMachine{}).
        WithEventFilter(predicates.ResourceNotPaused(log)).
        WithEventFilter(predicates.ResourceIsNotExternallyManaged(log)).
        Watches(
            &source.Kind{Type: &clusterv1.Machine{}},
            handler.EnqueueRequestsFromMapFunc(
                util.MachineToInfrastructureMapFunc(infrav1.GroupVersion.WithKind("PodmanMachine")),
            ),
        ) //TODO: add additional watches for PodmanCluster and Cluster if needed

    return builder.Complete(r)
}
```

# Implement controllers for your API types (2/2)

For **Reconcile** the following pattern is generally used:

1. Get the instance of the API type being reconciled
2. Get the owning CAPI type (i.e. if we reconciling **PodmanMachine** then get **Machine**)
3. If we don't have the machine then exit (owner reference isn't set yet)
4. Optionally, get the owning **Cluster** and Infra Cluster (i.e. **PodManCluster**)
5. If instance has a deletion timestamp, then in **reconcileDelete**:
  - a. do any actions to delete
  - b. remove finalizer and save
6. If instance has no deletion timestamp, then in **reconcileNormal**:
  - a. Add finalizer to instance and save
  - b. do any actions to create OR update

# Owner Reference

## When building a provider you should set ownerReference

- A link to a resource that is the owner
  - Example: Deployment owns Pods
  - Example: Cluster owns PodmanCluster
- Used heavily in Cluster API
- Implemented via the **metadata.ownerReference** field
- If the owner is deleted then either:
  - Cascading deletion (controlled via policy)
  - Orphaned resources

# Webhooks

If you need custom logic for defaults or validation you can create webhooks:

```
→ kubebuilder create webhook --group infrastructure --version v1alpha1 --kind PodmanCluster  
--defaulting --programmatic-validation  
→ kubebuilder create webhook --group infrastructure --version v1alpha1 --kind PodmanMachine  
--defaulting --programmatic-validation  
kubebuilder create webhook --group infrastructure --version v1alpha1 --kind PodmanMachineTemplate  
--defaulting --programmatic-validation
```

- Webhooks are Kubernetes Admission controllers
- This will scaffold both a **defaulting** & **validating** webhook
  - It's your responsibility to fill in the logic
- Defaulting webhook should only be used where kubebuilder defaults are not sufficient

# Webhook - Implementation



KubeCon



CloudNativeCon

Europe 2022

```
var _ webhook.Validator = &PodmanMachine{}

// ValidateCreate implements webhook.Validator so a webhook will be registered for the type
func (r *PodmanMachine) ValidateCreate() error {
    podmanmachinelog.Info("validate create", "name", r.Name)

    var allErrs field.ErrorList

    for _, mount := range r.Spec.ExtraMounts {
        if mount.HostPath == "" || mount.ContainerPath == "" {
            allErrs = append(allErrs, field.Invalid(
                field.NewPath("spec", "extraMounts"), "", "must specify both host and container path",
            ),
            )
        }
    }

    if len(allErrs) == 0 {
        return nil
    }

    return apierrors.NewInvalid(GroupVersion.WithKind("Cluster").GroupKind(), r.Name, allErrs)
}
```



# Local testing / development (1/2)

- Developing and debugging operators in Kubernetes can be painful.
- Tilt will save you a lot of time, pain and tears!
  - We need to tell Tilt about our provider via the **tilt-provider.json** file in repo root

```
[
  {
    "name": "podman",
    "config": {
      "image": "ghcr.io/capi-samples/cluster-api-provider-podman:dev",
      "live_reload_deps": [
        "main.go",
        "go.mod",
        "go.sum",
        "api",
        "controllers",
        "pkg"
      ],
      "label": "CAPP0D"
    }
  }
]
```

# Local testing / development (2/2)

- We can then follow the instructions from the CAPI docs to configure Tilt:
  - <https://cluster-api.sigs.k8s.io/developer/tilt.html>

```
{
  "default_registry": "gcr.io/capi-samples",
  "provider_repos": ["../github.com/capi-samples/cluster-api-provider-podman"],
  "enable_providers": ["podman", "kubeadm-bootstrap", "kubeadm-control-plane"],
  "kustomize_substitutions": {
    "EXP_CLUSTER_RESOURCE_SET": "true",
  },
  "extra_args": {
    "podman": ["--v=4"],
    "kubeadm-control-plane": ["--v=4"],
    "kubeadm-bootstrap": ["--v=4"],
    "core": ["--v=4"]
  },
  "debug": {
    "podman": {
      "continue": true,
      "port": 30000
    }
  }
}
```

# Testing



KubeCon



CloudNativeCon

Europe 2022

- Unit and integration tests are up to the provider implementers to follow the frameworks / strategy of their choice
- Envtest can be used for unit and integration tests
  - <https://github.com/kubernetes-sigs/controller-runtime/tree/master/pkg/envtest>
- With envtest, it is possible to interact with your provider like a real cluster
  - You can create / update CRDs and controllers can take action on events

```
By("bootstrapping test environment")
testEnv = &envtest.Environment{
    CRDDirectoryPaths: []string{
        filepath.Join("../", "../", "config", "crd", "bases"),
        filepath.Join(build.Default.GOPATH, "pkg", "mod", "sigs.k8s.io", "cluster-api@v1.1.3", "config", "crd", "bases"),
        filepath.Join(build.Default.GOPATH, "pkg", "mod", "sigs.k8s.io", "cluster-api@v1.1.3", "bootstrap", "kubeadm", "config", "crd", "bases")
    },
    ErrorIfCRDPathMissing: true,
}

var err error
cfg, err = testEnv.Start()
Expect(err).NotTo(HaveOccurred())
Expect(cfg).NotTo(BeNil())

err = infrastructurev1beta1.AddToScheme(scheme.Scheme)
Expect(err).NotTo(HaveOccurred())

err = clusterv1.AddToScheme(scheme.Scheme)
Expect(err).NotTo(HaveOccurred())

err = bootstrapv1.AddToScheme(scheme.Scheme)
Expect(err).NotTo(HaveOccurred())
```

# Testing Continued



KubeCon



CloudNativeCon

Europe 2022

- CAPI provides e2e framework - most of the code is reusable

```
It("Should create a workload cluster", func() {
    By("Creating a workload cluster")

    flavor := clusterctl.DefaultFlavor
    if input.Flavor != nil {
        flavor = *input.Flavor
    }

    clusterctl.ApplyClusterTemplateAndWait(ctx, clusterctl.ApplyClusterTemplateAndWaitInput{
        ClusterProxy: input.BootstrapClusterProxy,
        ConfigCluster: clusterctl.ConfigClusterInput{
            LogFolder:      filepath.Join(input.ArtifactFolder, "clusters", input.BootstrapClusterProxy.GetName()),
            ClusterctlConfigPath: input.ClusterctlConfigPath,
            KubeconfigPath:   input.BootstrapClusterProxy.GetKubeconfigPath(),
            InfrastructureProvider: clusterctl.DefaultInfrastructureProvider,
            Flavor:          flavor,
            Namespace:       namespace.Name,
            ClusterName:      fmt.Sprintf("%s-%s", specName, util.RandomString(6)),
            KubernetesVersion: input.E2EConfig.GetVariable(KubernetesVersion),
            ControlPlaneMachineCount: pointer.Int64Ptr(1),
            WorkerMachineCount: pointer.Int64Ptr(1),
        },
        WaitForClusterIntervals: input.E2EConfig.GetIntervals(specName, "wait-cluster"),
        WaitForControlPlaneIntervals: input.E2EConfig.GetIntervals(specName, "wait-control-plane"),
        WaitForMachineDeployments: input.E2EConfig.GetIntervals(specName, "wait-worker-nodes"),
    }, clusterResources)

    By("PASSED!")
})
```

# Releasing

To be installable via **clusterctl init** you must:

- Publish your provider as a container to a registry
- Create a GitHub release:
  - Release name should be a version number following the semver convention
  - Attach the following assets:
    - metadata.yaml
    - infrastructure-components.yaml
    - cluster-template\*.yaml

How do I generate Infrastructure-components.yaml?

```
kustomize build config/default/ > infrastructure-components.yaml  
# NOTE: replace container image with the one from the registry
```

# Community

- Building a provider is just the start
- It's advisable to get involved in the wider CAPI community
  - Attend the office hours calls on Wednesdays
  - Read & comment on issues and enhancement proposals (CAEP)
  - Update your provider when new CAPI versions are released
- To raise awareness or to increase adoption for your new provider
  - Host regular Office Hours
  - Encourage new contributors by having a well-defined README, good first issues
  - Use forums like CAPI Office Hours to talk about your provider
- To donate to kubernetes-sigs
  - Check if your repo follows the [kubernetes template project](#) format
  - Fill out the repo [migration request](#)
  - Stay on top of the request and answer any queries :)

# Wrapping up

There is a lot we haven't covered. Some important areas:

- Multiple API versions, conversions and Hub/Spoke

Some resources when you implement your own provider:

- CAPI Repo: <https://github.com/kubernetes-sigs/cluster-api>
- CAPI Provider Implementers docs: <https://cluster-api.sigs.k8s.io/developer/providers/implementers.html>
- List of existing providers: <https://cluster-api.sigs.k8s.io/reference/providers.html>
- Kubebuilder docs: <https://book.kubebuilder.io/>
- Sample Podman provider: <https://github.com/capi-samples/cluster-api-provider-podman>



KubeCon



CloudNativeCon

Europe 2022

Thanks for listening.....any questions?

