



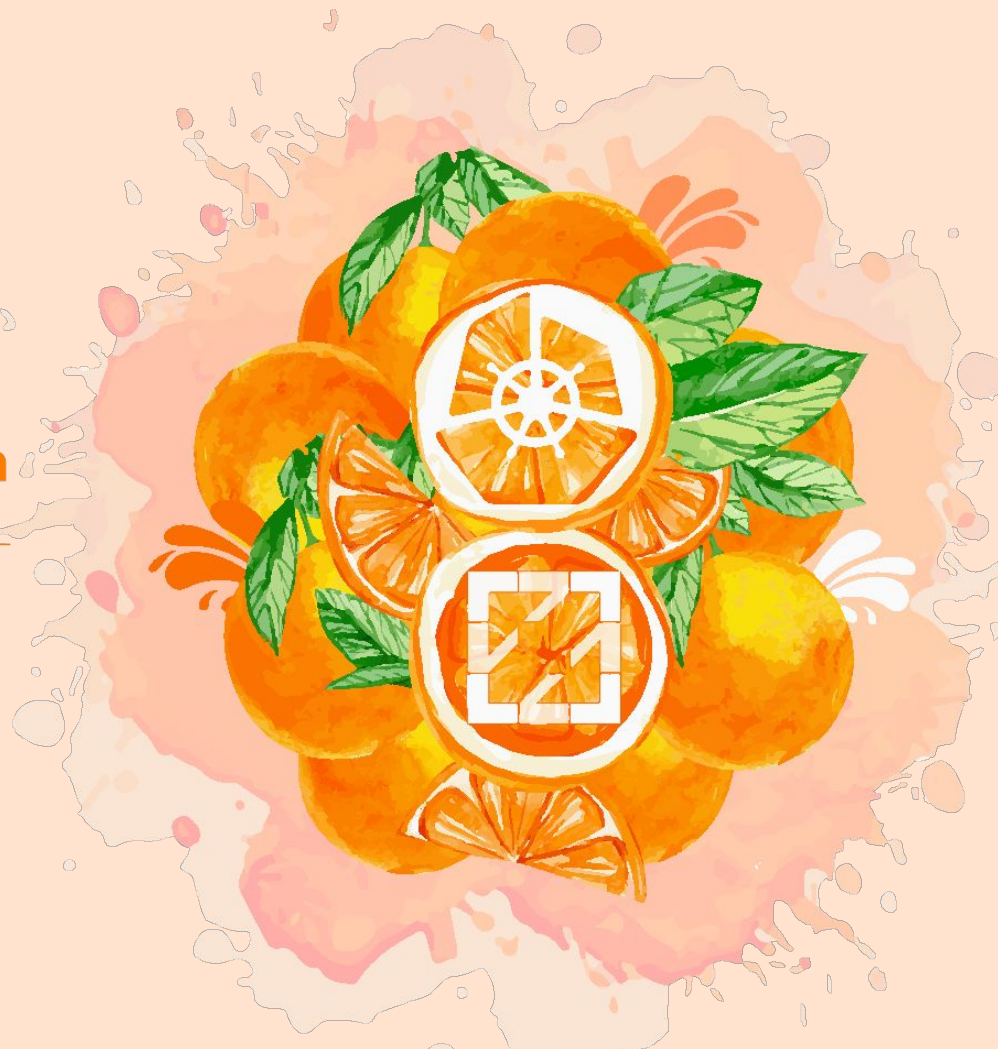
KubeCon



CloudNativeCon

Europe 2022

WELCOME TO VALENCIA





KubeCon



CloudNativeCon


Europe 2022

Adapting TiKV for Cloud Storage

Xinye Tao ([@tabokie](#)), PingCAP




- A **distributed** transactional key-value storage engine
 - Scale out to hundreds of nodes
 - Replication of both WAL and data files



PromCon
North America 2021

TiKV + Cloud Storage



PromCon
North America 2021


- A **distributed** storage engine for TiDB
 - Scale out to hundreds of nodes
 - Replication of both WAL and data files
- **Cloud-based** disk: Elastic Block Storage (AWS), Persistent Disk (GCP), Managed Disk (Azure)
 - Under the hood, they are connected to remote shared hardware and replicated across different zones
 - High latency
 - Provisioned performance (bandwidth, IOPS)
 - Service degradation and outage

Challenges



PromCon
North America 2021

- Scalability
 - Likelihood of rare events increases on a larger scale
 - Storage performance is more likely to degrade on a larger scale
- Cost
 - User are more sensitive to read/write amplification
 - There is an additional cost to replicating data flow




- Storage Performance Budgets (BW, IOPS)
- Large Events (Analyze, Backup)
- Background Writes (Compaction, GC)
- User Writes

Challenges



PromCon
North America 2021

- Scalability
 - Likelihood of rare events increases on a larger scale
 - Storage performance is more likely to degrade on a larger scale
- Cost
 - User are more sensitive to space amplification and write amplification
 - There is an additional cost to replicating data flow



Raft Engine



- A lightweight log store in Rust (<https://github.com/tikv/raft-engine>)
- [x] <Primary Goal> write less than RocksDB, both foreground and background
- [] <Secondary Goal> a more performant engine than RocksDB

Raft Engine \approx Index + Logs



- A lightweight log store in Rust (<https://github.com/tikv/raft-engine>)
- [x] <Primary Goal> write less, both foreground and background
 - An **in-memory index** keeps track of all active log entries
 - No need to sort the log files
 - No need for GC to read deleted data
 - Log entries are **compressed** and appended to log files
 - Test results* showed **30%** reduction in server write I/Os

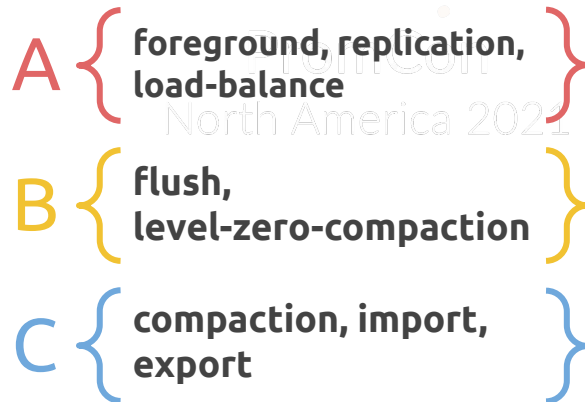
PromCon
North America 2021

** Ran TPC-C 50 threads on a 3 node cluster, AWS r5.2xlarge with gp3 disk*

Priority I/O Scheduling



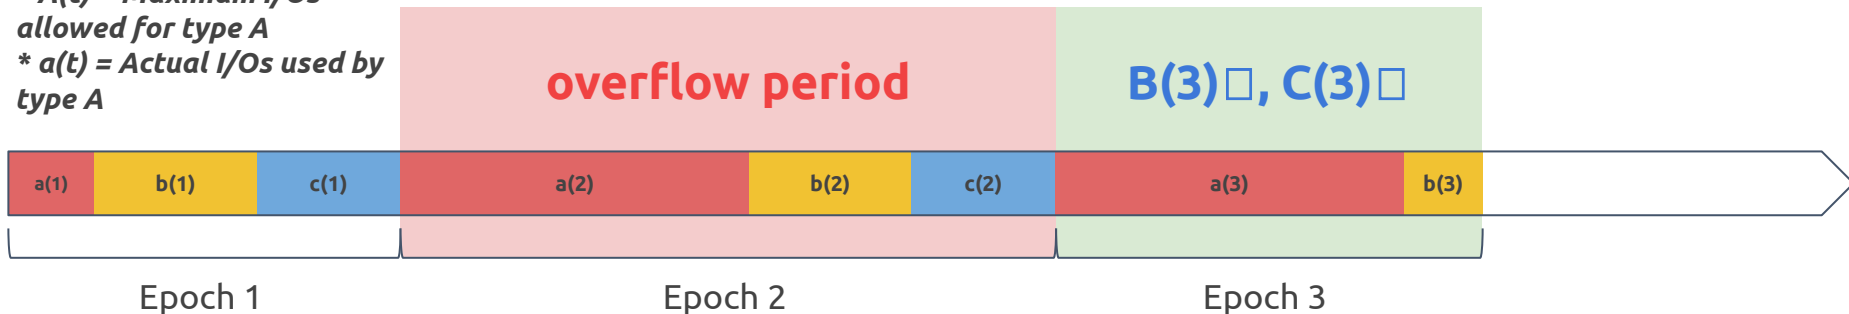
- A non-intrusive way to prioritized I/O requests ([#9197](#))
 - No userland I/O queue, zero overhead
- Read I/O accounting and Rust async support ([#11969](#))
 - Pulling stats from ``proc/task/io`` when needed



Notes:

* $A(t)$ = Maximum I/Os allowed for type A

* $a(t)$ = Actual I/Os used by type A




Priority I/O Scheduling



PromCon
North America 2021

- A non-intrusive way to prioritized I/O requests ([#9197](#))
- I/O intensive tasks are well constrained



*Importing data while running TPC-C workload
w/o I/O scheduling*

More on the way



- CPU Limiting: proactive back pressure for low-resource environments ([#12151](#))
- Raft Witness: a write-only node that only replicates logs ([raft-rs#145](#))
- ...