



Jaeger present and future

Maintainer Talk


Kubecon EU 2022

Jonah Kowall [@jkowall]

Pavol Loffay [@ploffay]



Jonah Kowall?



 jkowall



Pavol Loffay?

 ploffay





Agenda

1. Intro to Distributed Tracing and Jaeger (Jonah)
2. Jaeger and OpenTelemetry (Pavol)
3. New Monitoring tab and Prometheus support (Jonah)
4. Jaeger Kubernetes Operator deployment (Pavol)
5. New Key Features + Roadmap for Jaeger (Jonah)
6. Q&A from audience in room and online (Jonah + Pavol)

Intro to Distributed Tracing and Jaeger



OpenTelemetry Semantics

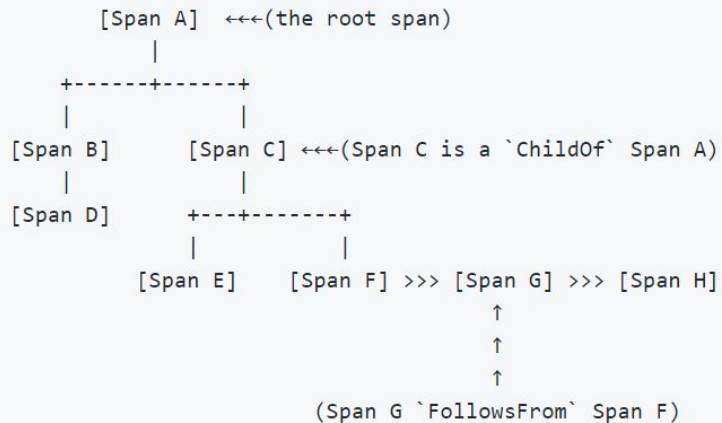
Trace represents an end-to-end request (and response); made up of single or multiple **Spans**

Span represents work done by a single-service or component with time intervals and associated metadata such as **Tags**

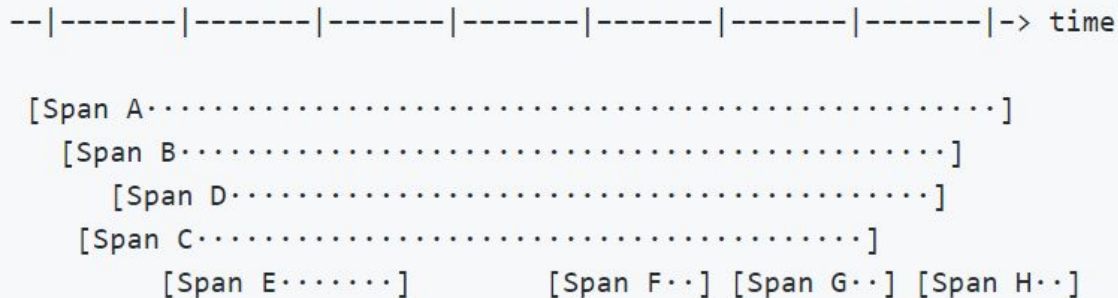
Tags contain metadata to help contextualize a span

Relationships in tracing

Causal relationships between Spans in a single Trace



Temporal relationships between Spans in a single Trace




Measure errors, latency, and other indicators across each span

← ▼ frontend: HTTP GET /dispatch 4737e2c

Search...



Trace Start **December 16, 2018 5:19 PM** Duration **700.68ms** Services **6** Depth **5** Total Spans **50**



Service & Operation



0ms 175.17ms 350.34ms 525.51ms 700.68ms

▼ frontend HTTP GET /dispatch

▼ frontend HTTP GET: /customer

▼ frontend HTTP GET

▼ customer HTTP GET /customer

mysql SQL SELECT

▼ frontend Driver::findNearest

▼ driver Driver::findNearest

redis FindDriverIDs

redis GetDriver

❌ redis GetDriver

redis GetDriver

redis GetDriver

redis GetDriver

redis GetDriver

redis GetDriver

redis GetDriver

redis GetDriver

❌ redis GetDriver

redis GetDriver

redis GetDriver

199.4ms

199.08ms

18.81ms

10.55ms

34.14ms

9.27ms

9.83ms

9.84ms

10.25ms

15.86ms

13.55ms

12.29ms

36.42ms

7.35ms

10.37ms

Jaeger and OpenTelemetry



OpenTelemetry and Jaeger

- Jaeger - platform
- OpenTelemetry - data collection

OpenTelemetry Components



REFERENCE ARCHITECTURE

- C++
- › .NET
- › Erlang/Elixir
- › Go
- › Java
- › JavaScript
- › PHP
- › Python
- › Ruby
- Rust
- Swift

OpenTelemetry

DocsRegistryStatusCommunityBlog

Registry

Find libraries, plugins, integrations, and other useful tools for extending OpenTelemetry.


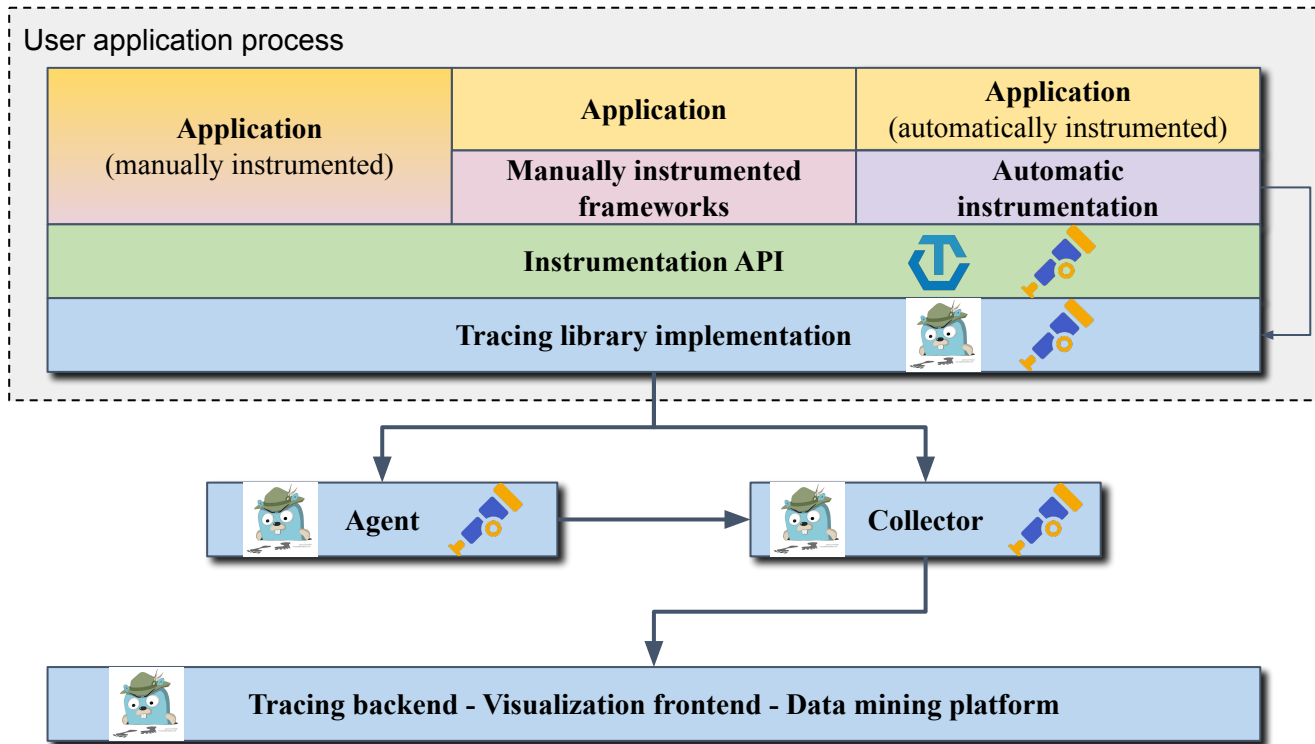
What do you need?

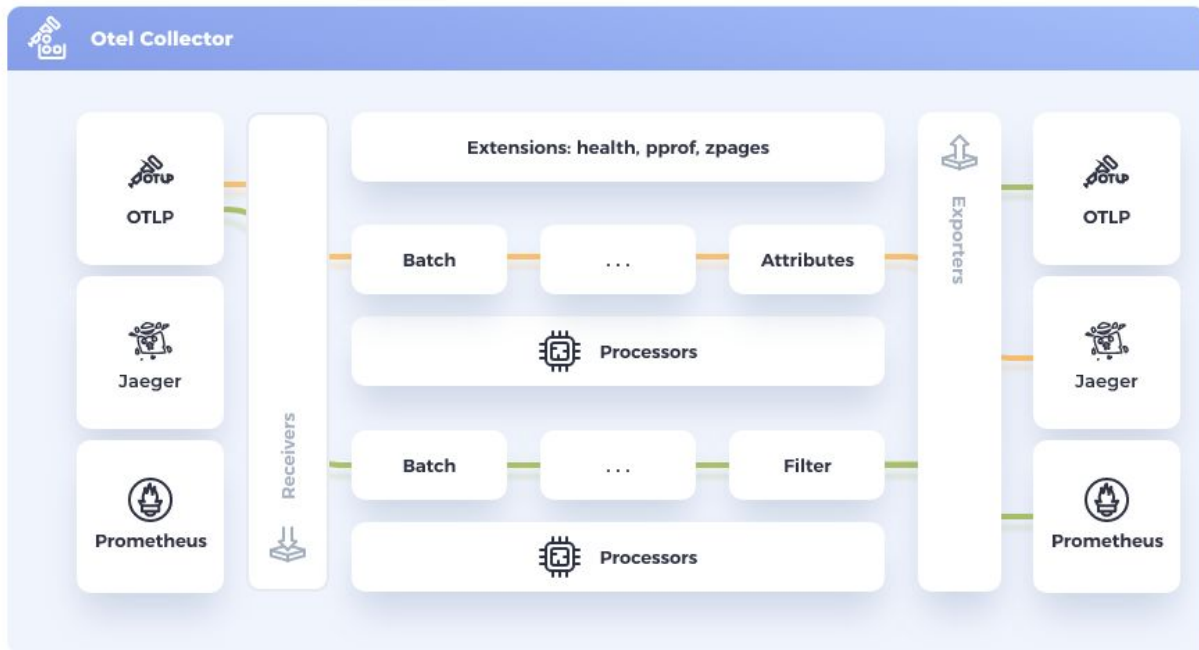
The OpenTelemetry Registry allows you to search for instrumentation libraries, tracer implementations, utilities, and other useful projects in the OpenTelemetry ecosystem.

- Not able to find an exporter for your language? Remember, the [OpenTelemetry Collector](#) supports exporting to a variety of systems and works with all OpenTelemetry Core Components!
- Are you a project maintainer? See, [Adding a project to the OpenTelemetry Registry](#).
- Check back regularly, the community and registry are growing!

[.NET](#)
The OpenTelemetry API and SDK for .NET (C#, F#)
[ActionPack Instrumentation](#)
ActionPack instrumentation for Ruby.
[ActionView Instrumentation](#)
ActionView instrumentation for Ruby.
[Active Model Serializers Instrumentation](#)
Active Model Serializers instrumentation for Ruby.

[dotnet](#)[core](#)
[ruby](#)[instrumentation](#)
[ruby](#)[instrumentation](#)
[ruby](#)[instrumentation](#)





OTEL COLLECTOR



Instrumentation

- Jaeger clients are deprecated in favor of OpenTelemetry
- OpenTelemetry SDKs support
 - Jaeger context-propagation header
 - Jaeger remote sampler



OpenTelemetry Collector

- Jaeger receiver
 - proto over gRPC (default endpoint = 0.0.0.0:14250)
 - thrift_binary (default endpoint = 0.0.0.0:6832)
 - thrift_compact (default endpoint = 0.0.0.0:6831)
 - thrift_http (default endpoint = 0.0.0.0:14268)
- Jaeger exporter
 - proto over gRPC
- Jaeger remote sampler extension
 - serves HTTP
- Kafka receiver/exporter
 - Jaeger proto
 - Jaeger JSON



Jaeger V3 Query API

```
// Response object with spans.
message SpansResponseChunk {
  // A list of OpenTelemetry ResourceSpans.
  // In case of JSON format the ids (trace_id, span_id, parent_id)
  //   are encoded in base64 even though OpenTelemetry specification
  //   mandates to use hex encoding [2].
  // Base64 is chosen to keep compatibility with JSONPb codec.
  // [1]:
https://github.com/open-telemetry/opentelemetry-proto/blob/main/opentelemetry/proto/trace/v1/trace.proto
  // [2]:
https://github.com/open-telemetry/opentelemetry-specification/blob/main/specification/protocol/otlp.md#otlphttp

  repeated opentelemetry.proto.trace.v1.ResourceSpans resource_spans = 1;
}
```

New Monitoring tab and Prometheus support

Tracing and Monitoring



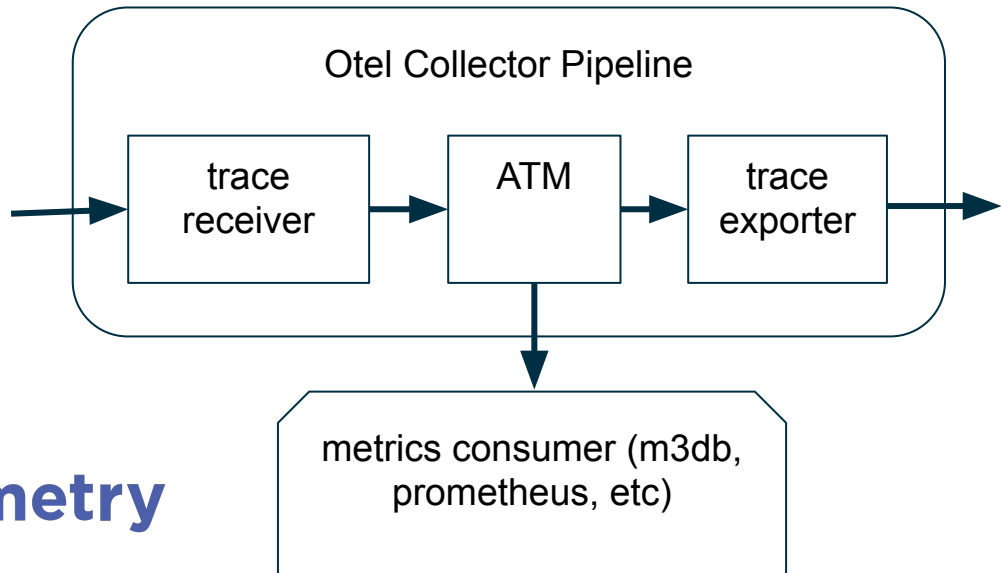
- What's the difference between “distributed tracing” and “Application Performance Monitoring - APM”?
 - Traces / Events
 - Metrics
- Use cases
 - Monitoring
 - Alerting
 - Planning

Aggregated Trace Metrics (ATM)



- Prometheus can handle all of these metrics use cases
- We just have to generate the metrics from the traces

Derive aggregated metrics from traces



SpanMetrics Processor in OpenTelemetry

```
processors:
  batch:
  spanmetrics:
    metrics_exporter: otlp/spanmetrics
    latency_histogram_buckets: [100us, 1ms, 2ms, 6ms, 10ms, 100ms, 250ms]
    dimensions:
      - name: http.method
        default: GET
      - name: http.status_code

service:
  pipelines:
    traces:
      receivers: [jaeger]
      processors: [spanmetrics, batch]
      exporters: [jaeger]

    # The exporter name must match the metrics_exporter name.
    # The receiver is just a dummy and never used; added to pass validation requiring at least one receiver in a pipeline.
    metrics/spanmetrics:
      receivers: [otlp/spanmetrics]
      exporters: [otlp/spanmetrics]

  metrics:
    receivers: [otlp]
    exporters: [prometheus]
```

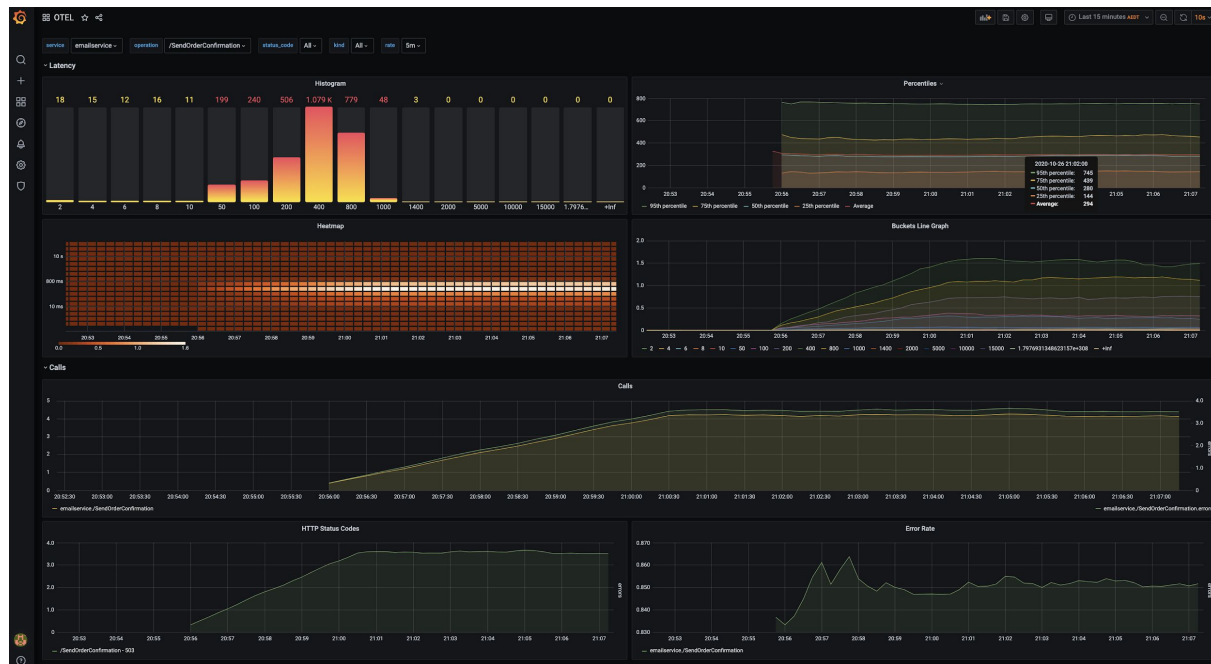
Define how many metrics (in this case method type and status code) and the buckets.

Result: Generate a metric per bucket per status code

Using Prometheus Metrics

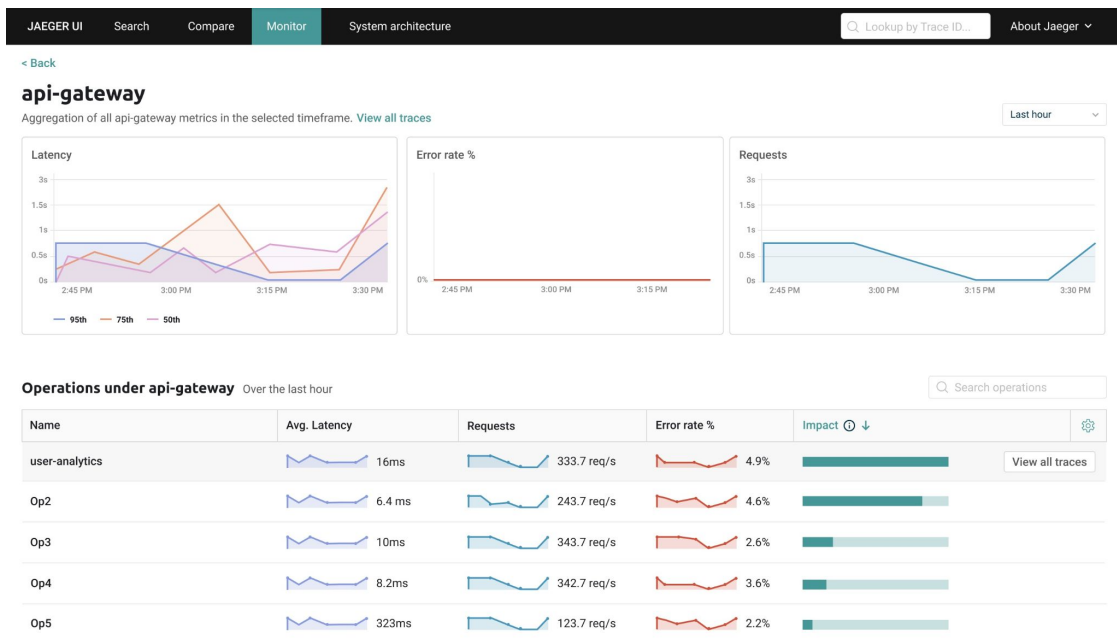
MetricQuery service in Jaeger to query metric backends.

- First support will be for promql compatible backends (ex: Prometheus, Cortex, Thanos, M3DB)
- Community can add other systems as needed



Using Prometheus Metrics inside Jaeger UI

- New “monitor” homepage in Jaeger to provide status and health of transactions



Jaeger Kubernetes Operator



Jaeger Operator



“operates” Jaeger on Kubernetes

- github.com/jaegertracing/jaeger-operator
- jaegertracing.io/docs/latest/operator

Getting Started

- jaegertracing.io/docs/1.34/operator
- Storage schema creation
- Jaeger upgrades
- Auto recognizes available APIs - OpenShift/Kubernetes, ES/Strimzi
- Can generate plain Kubernetes manifest files



Jaeger CRD



- CRD - Custom Resource Definition

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  strategy: allInOne | production | streaming
```

- `kubectl get jaegers`




Jaeger Operator


Strategies



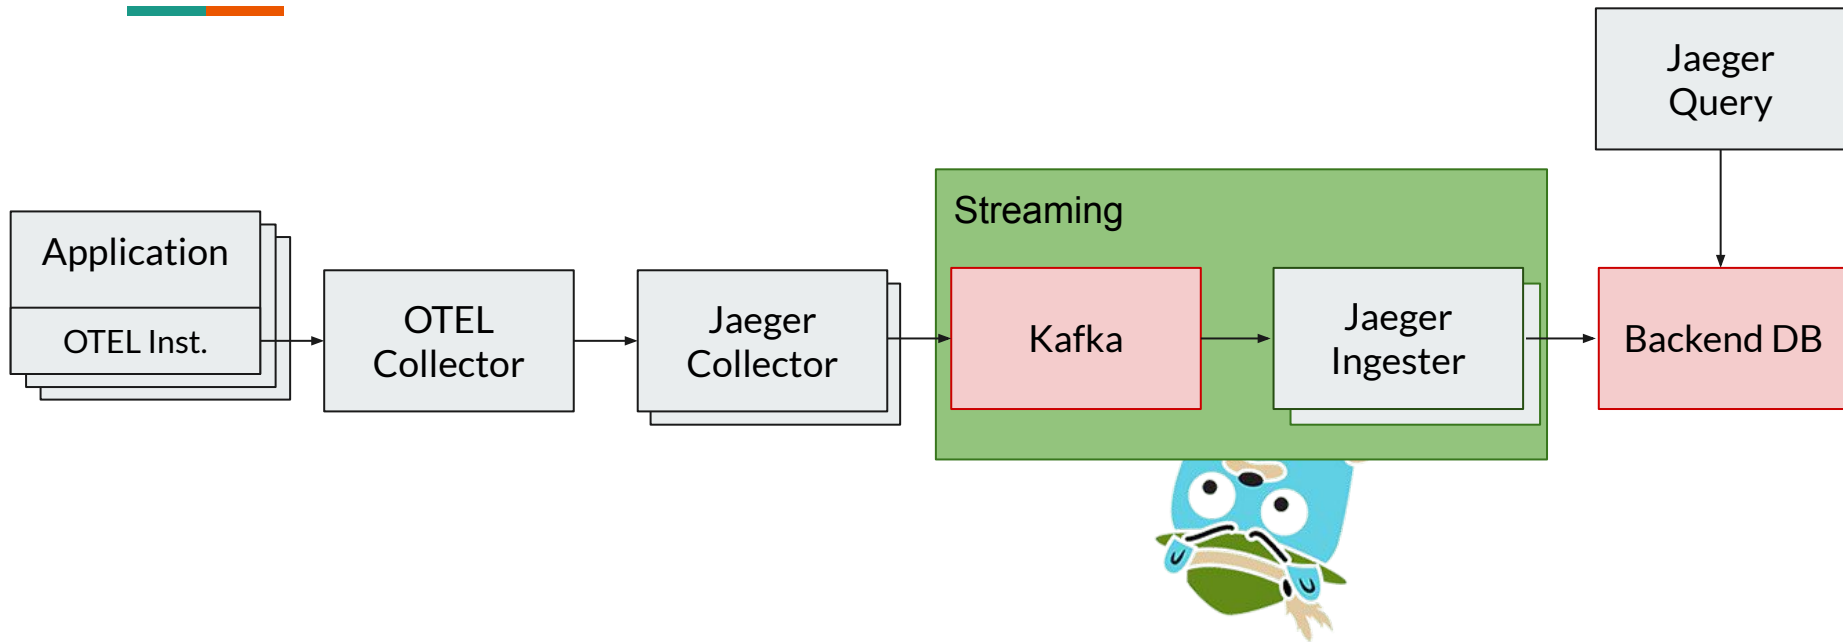
Strategies: all-in-one



Strategies: production



Strategies: streaming





Jaeger Operator



Storage configuration

Storage Configuration (backend)



spec.storage

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  storage:
    type: memory|elasticsearch|cassandra|badger|grpc-plugin|kafka
    options:
      es:
        server-urls: http://some-elastic-cluster-somewhere:9200
```


Detailed Configuration

```
$ docker run -e SPAN_STORAGE_TYPE=elasticsearch jaegertracing/jaeger-collector --help
...
--collector.num-workers int      The number of workers pulling items from the queue
--collector.queue-size int       The queue size of the collector
--collector.queue-size-memory uint (experimental) The max memory size in MiB
...
```

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  collector:
    options:
      collector:
        queue-size: 100
```





Jaeger Operator

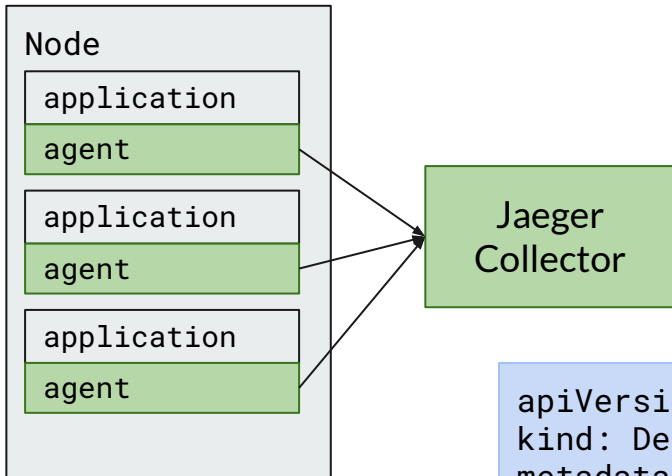


Agent strategies

Agent Strategy (Sidecar)




Sidecar



```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-dev
spec:
  agent:
    strategy: Sidecar
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
  annotations:
    "sidecar.jaegertracing.io/inject": "true"
```

Agent Strategy (DaemonSet)



```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  agent:
    strategy: DaemonSet
```



Remote Sampling

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  sampling:
    options:
      default_strategy:
        type: probabilistic
        param: 0.5
      service_strategies:
        - service: foo
          type: probabilistic
          param: 0.8
        operation_strategies:
          - operation: get
            type: probabilistic
            param: 0.2
```

Autoscaling Collectors and Ingesters

- Collectors and Ingesters
- By default creates an HPA with max of 100

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  collector:
    maxReplicas: 20
```

```
apiVersion: jaegertracing.io/v1
kind: Jaeger
metadata:
  name: jaeger-cluster
spec:
  collector:
    autoScale: false
    replicas: 10
```



Operator Integrations



- Storage
 - Kafka - Strimzi
 - Elasticsearch - OpenShift cluster logging
- Monitoring Jaeger operator
 - Prometheus
 - OpenTelemetry

New Key Features + Roadmap

New Key Features



- Adaptive Sampling - Jaeger backend can be configured to perform fully automated and dynamic control of sampling rates based on predefined targets
- Service Performance Monitoring - We covered this
- All-in on OpenTelemetry - The Jaeger Clients/SDKs have been officially retired in favor of OpenTelemetry

Roadmap



- Updates to dependency graphs
 - Normalize the three types of graphs in Jaeger
 - Overlay service performance metrics on graph
 - Potentially move calculations from Spark/Kafka streams to OpenTelemetry collector
- Move towards OpenTelemetry collector
 - Remove the need for Jaeger collector and normalize on a distribution of the collector for writes to Jaeger data stores
 - Native OTLP support

And more interesting capabilities coming in the future

Q&A from audience in room and online

Resources



jaegertracing.io/docs



[monthly community call and Notes](#)
[CNCF Slack #jaeger : https://slack.cncf.io](https://slack.cncf.io)



[@jaegertracing](https://twitter.com/jaegertracing)



medium.com/jaegertracing

