



**KubeCon**



**CloudNativeCon**

**Europe 2022**

**WELCOME TO VALENCIA**





KubeCon



CloudNativeCon

Europe 2022

# Sailing Multi Cloud Traffic Management With Karmada

Zhonghu Xu, Huawei



# About me

- Istio steering committee member
- Istio Core Maintainer & Contributor
- Open source enthusiastic
- Kubernetes member & core contributor
- Github: <https://github.com/hzxuzhonghu>



**Zhonghu Xu**  
Open Source Engineer  
*Huawei*



KubeCon



CloudNativeCon

Europe 2022

PromCon  
North America 2021

# Agenda

1. Background
2. Karmada introduction
3. Multi cloud traffic management
4. Conclusion

# Background: Multi cloud, multi cluster strategy is widely adopted



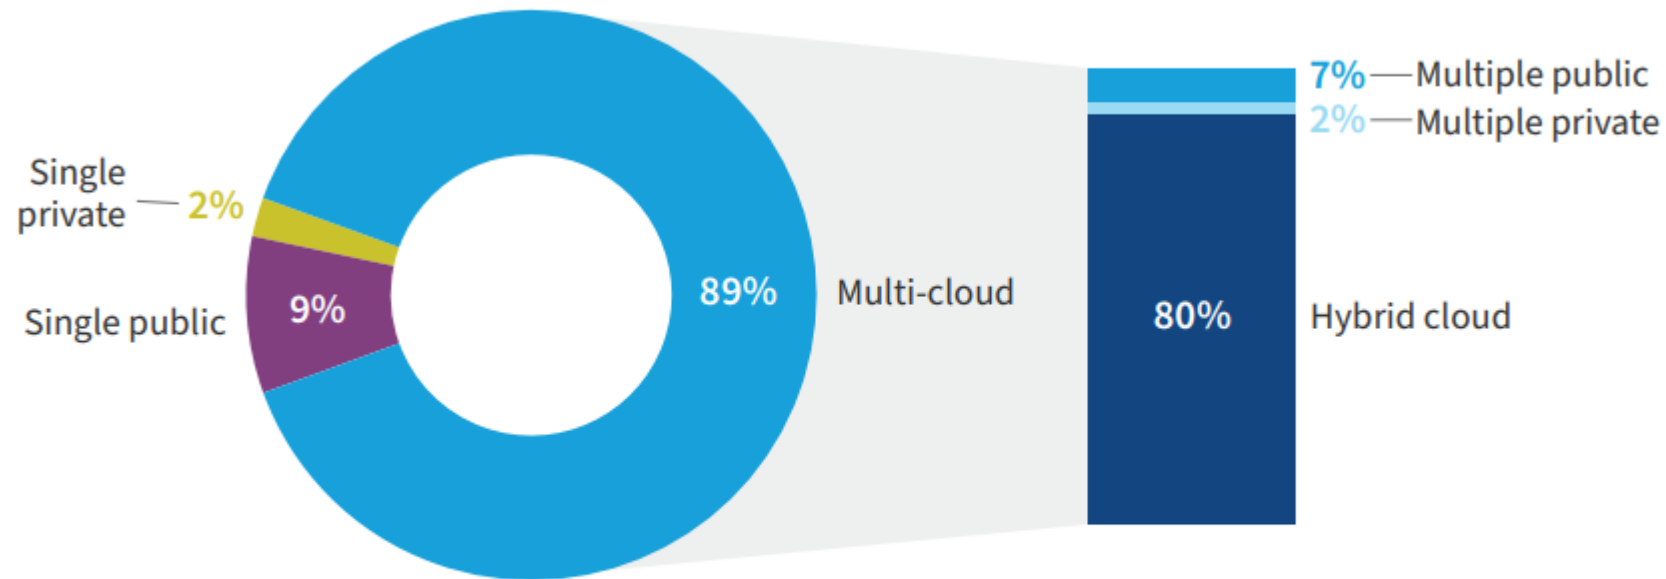
KubeCon



CloudNativeCon

Europe 2022

Cloud strategy for all organizations



N=753

Source: Flexera 2022 State of the Cloud Report

**flexera**

# Background: Why Multi Cloud

- Avoid vendor locking
- Meeting compliance requirements
- Enhancing resilience
- Improving flexibility and scalability

# Background : Multi Cloud Challenges

- Management complexity
- Security concerns.
- Communication
- Monitoring system concern

# Background: Multi-Cloud Kubernetes Challenge

## Challenges of managing multiple container clusters

---

### Too Many Clusters

Cumbersome and Repetitive setup  
Incompatible Cluster Lifecycle API  
Fragmented API endpoints

### Workload Fragmentation

Per-cluster customization for Apps  
Multi-cluster service discovery for Apps  
Sync Apps between clusters

### Boundary of Clusters

Resource Scheduling  
Application Availability  
Horizontal Auto-scaling

### Vendor lock-in

Deployment Gravity  
Lack of Migration Automation  
Lack of independent, neutral, open source multi-cluster management projects



# Karmada: Open, Cloud-Native, Multi-Cloud Orchestration Engine



Easily build infinitely scalable cluster pools with Karmada  
Use multi-cloud clusters just like a single K8s cluster

## K8s Native API Compatible

Zero change upgrade: single-cluster → multi-cluster  
Seamless integration of existing K8s tool chain

## Open and Neutral

Jointly initiated by Internet, finance,  
manufacturing, teleco, cloud providers, etc.  
Target for open governance with CNCF

## Avoid Vendor Lock-in

Integration with mainstream cloud providers  
Automatic allocation, migration across clusters  
Not tied to proprietary vendor orchestration

## Out of the Box

Built-in policy sets for scenarios:  
GDPS, Active-active, Remote Disaster Recovery

## Fruitful Multi-Cluster Scheduling Policies

Cluster Affinity, Multi Cluster Splitting/Rebalancing,  
Multi-Dimension HA: Region/AZ/Cluster/Provider

## Centralized Management

Cluster location agnostic  
Support clusters in Public cloud, on-prem or edge

# Karmada Architecture

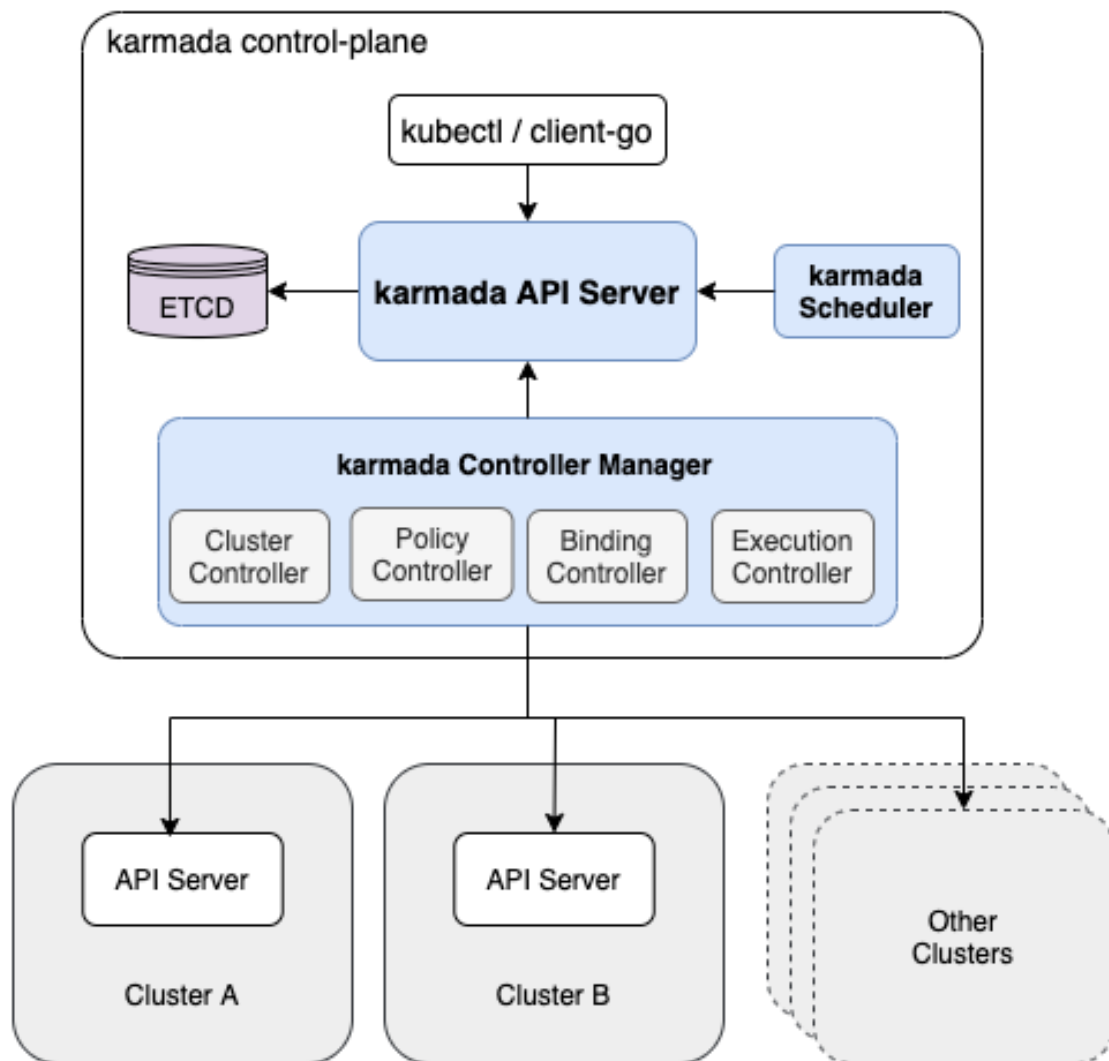


KubeCon

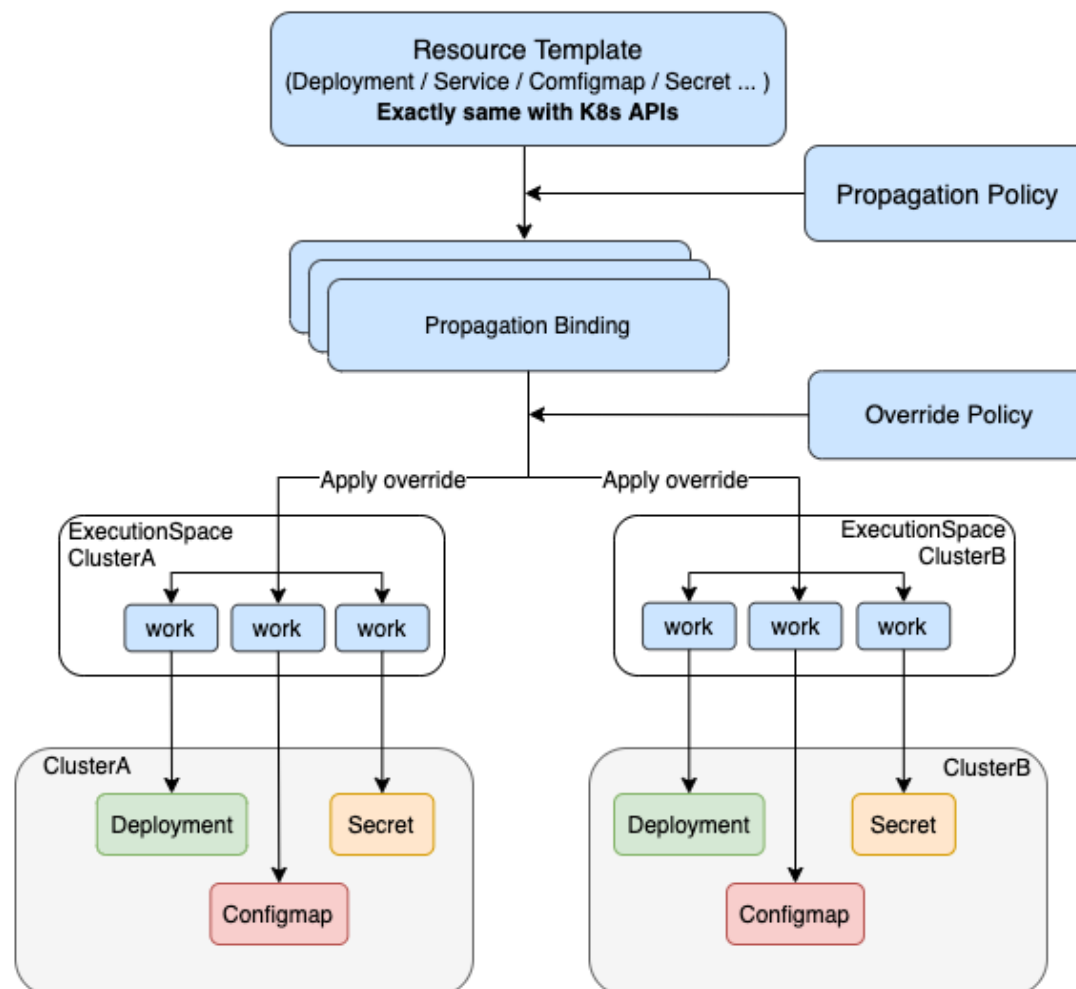


CloudNativeCon

Europe 2022



## Karmada Resources



# Karmada API Workflow

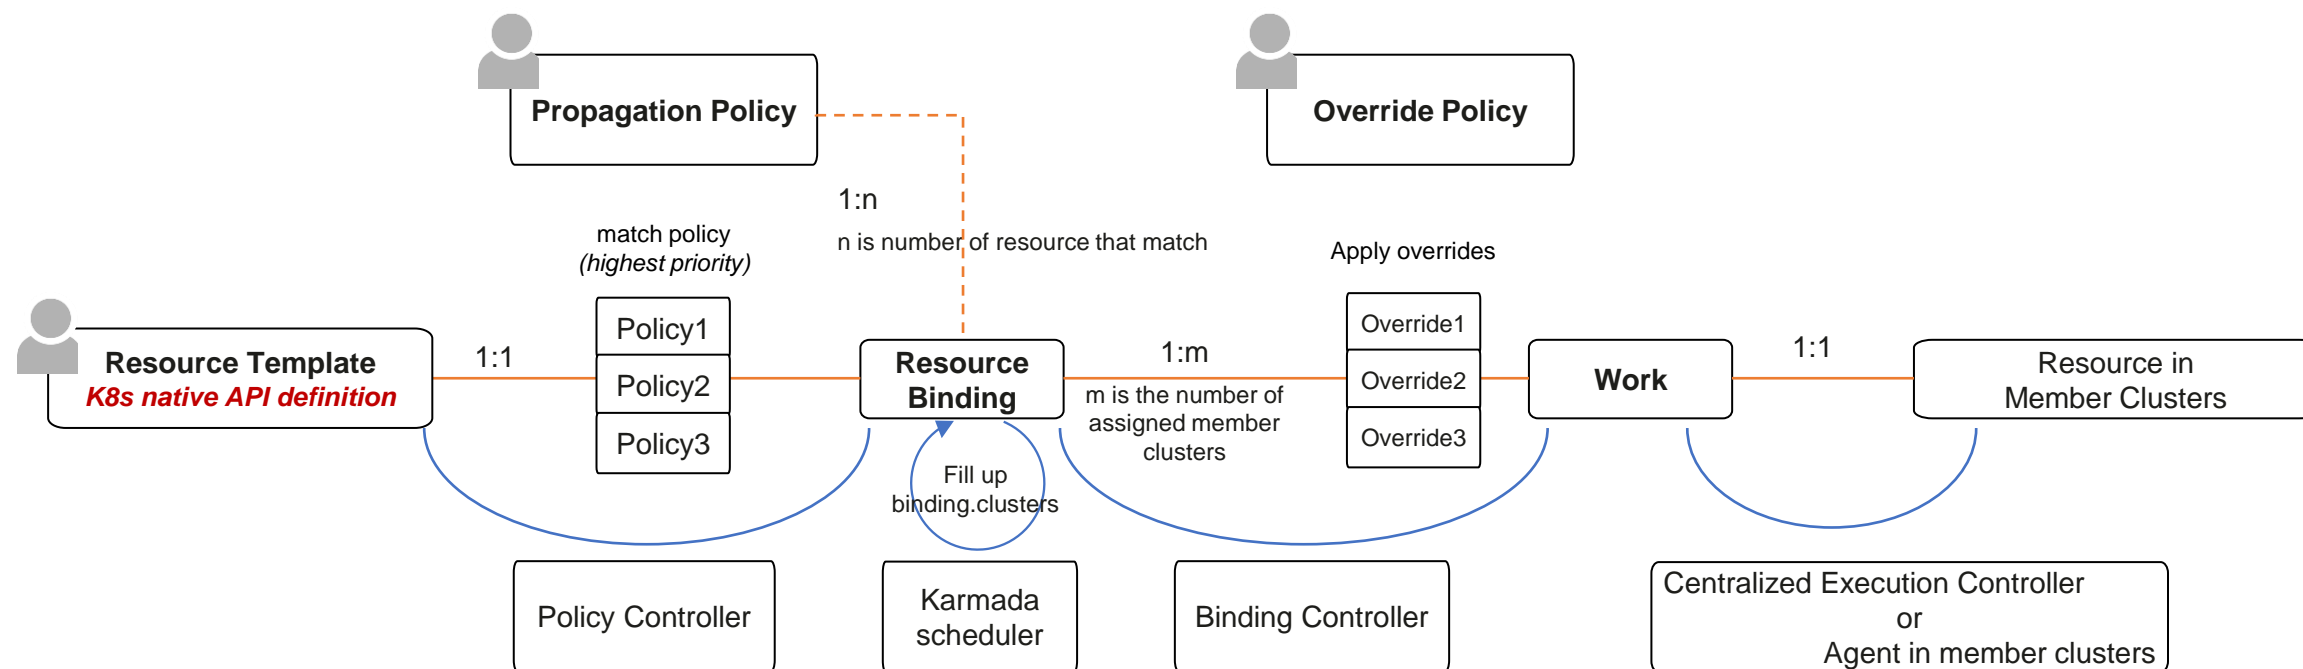


KubeCon



CloudNativeCon

Europe 2022



# Zero Change: Running Multi-Cluster Application with vanilla K8s API



KubeCon



CloudNativeCon

Europe 2022

## Reusable propagation policy

```
apiVersion: policy.karmada.io/v1alpha1
kind: PropagationPolicy
metadata:
  name: multi-zone-replication
spec:
  resourceSelectors:
  - apiVersion: apps/v1
    kind: Deployment
    labelSelector:
      matchLabels:
        ha-mode: multi-zone-replication
  placement:
    spreadConstraints:
    - spreadByField: zone
      maxGroups: 3
      minGroups: 3
```

Example: spread all deployments (that has label  
ha-mode: multi-zone-replication ) to 3 zones

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
    ha-mode: multi-zone-replication
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        ports:
        - containerPort: 80
```

Exactly vanilla  
Kubernetes API

Use exactly vanilla K8s API to deploy multi-cluster application  
kubectl create -f nginx-deployment.yaml

# Propagation Policy: multi-cluster scheduling policy



KubeCon



CloudNativeCon

Europe 2022

```
apiVersion: policy.karmada.io/v1alpha1
kind: PropagationPolicy
metadata:
  name: example-policy
spec:
  resourceSelectors:
    - apiVersion: apps/v1
      kind: Deployment
      name: deployment-1
      labelSelector: # standard labelSelector
  propagateDependencies: false
  placement:
    clusterAffinity:
      clusterNames:
        - cluster1
        - cluster3
    clusterTolerations: # like pod tolerations
  spreadConstraints:
    - spreadByField: zone
      maximum: 3
      minimum: 3
  schedulerName: default
```

## *resourceSelector*

- Match resources that the propagation policy apply to.
- *apiVersion + kind* for basic filtering
- *name* for exact match
- *labelSelector* for advanced matching

## *placement*

- Represents preferences of propagating resources
- *clusterAffinity*:
  - Preferred clusters to go
  - Exact match by names or match by labelselector
- *clusterTolerations*:
  - similar idea to pod tolerations and node taints, member clusters have taints to mark reservation for special usage, and only federated resources spread by propagation policies with corresponding tolerations can go there.
- *spreadConstraints*:
  - constraints of spreading federated resources among member clusters.
  - Users can specify dynamic grouping clusters by labels or by fields, and maximum/minimum groups they want to use.

# Override Policy: resource customization among clusters

```
apiVersion: policy.karmada.io/v1alpha1
kind: OverridePolicy
metadata:
  name: example-override
  namespace: default
spec:
  # restrict resource types that this override policy applies to
  resourceSelectors:
    - apiVersion: apps/v1
      kind: Deployment
      name: # user can either select resource by name or by labelselector
      labelSelector:
        matchLabels:
          image: nginx
  # this override policy will only apply to resources propagated to the matching clusters
  targetCluster:
    clusterNames: # user can either select cluster by names or by labelselector
      - dc-1-cluster-1
      - dc-1-cluster-2
    labelSelector:
      matchLabels:
        failedomain.kubernetes.io/region: dc1
  # all matching targetClusters would share the same set of overrides below
  overrides:
    plaintext:
      - path: "/spec/template/spec/containers/0/image"
        value: "dc-1.registry.io/nginx:1.17.0-alpine"
      - path: "/metadata/annotations"
        op: "add"
        value:
          foo: bar
      - path: "/metadata/annotations/foo"
        op: "remove"
```

## Example usage of override policy API

There are 3 clusters:

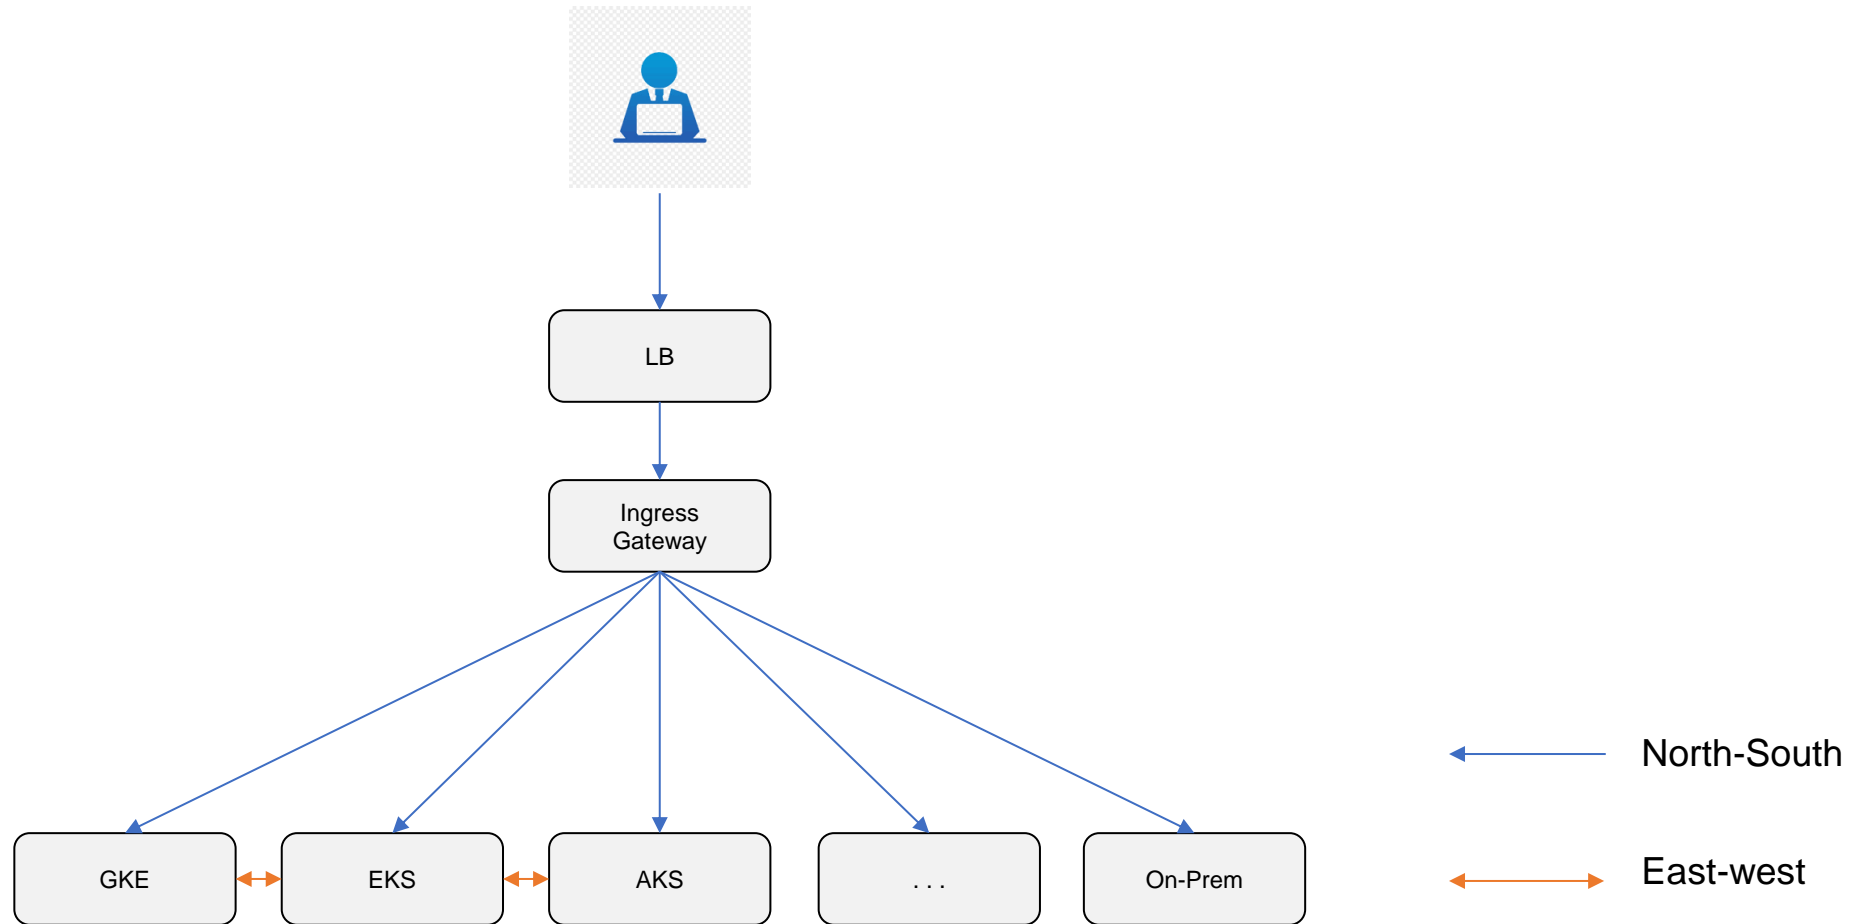
- *cluster-1* and *cluster-2* locate in *dc-1* which is an on-prem environment
- *cluster-3* is a managed Kubernetes cluster on public cloud.

To save image downloading bandwidth, latency:

- For deployments in *cluster-1* and *cluster-2*, download image from local registry *dc-1.registry.io*
- For deployments in *cluster-3*, download image from the registry managed by cloud provider

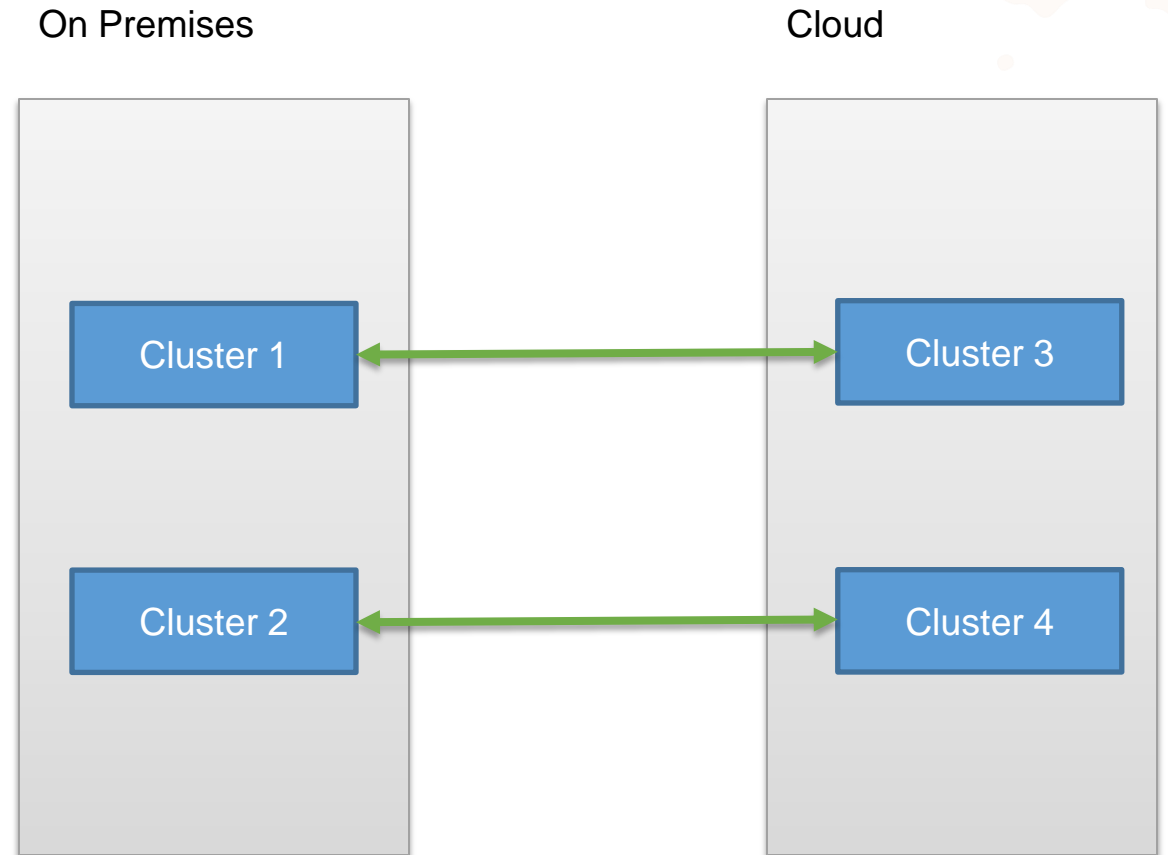
Override rules for *cluster-1* and *cluster-2* shown as in the left

# Multi cloud Traffic Management



# Multi cloud Traffic Management Challenges

- Network reachability
- Service discovery
- DNS resolve
- Load balancer Policy
- Security for cross cluster traffic

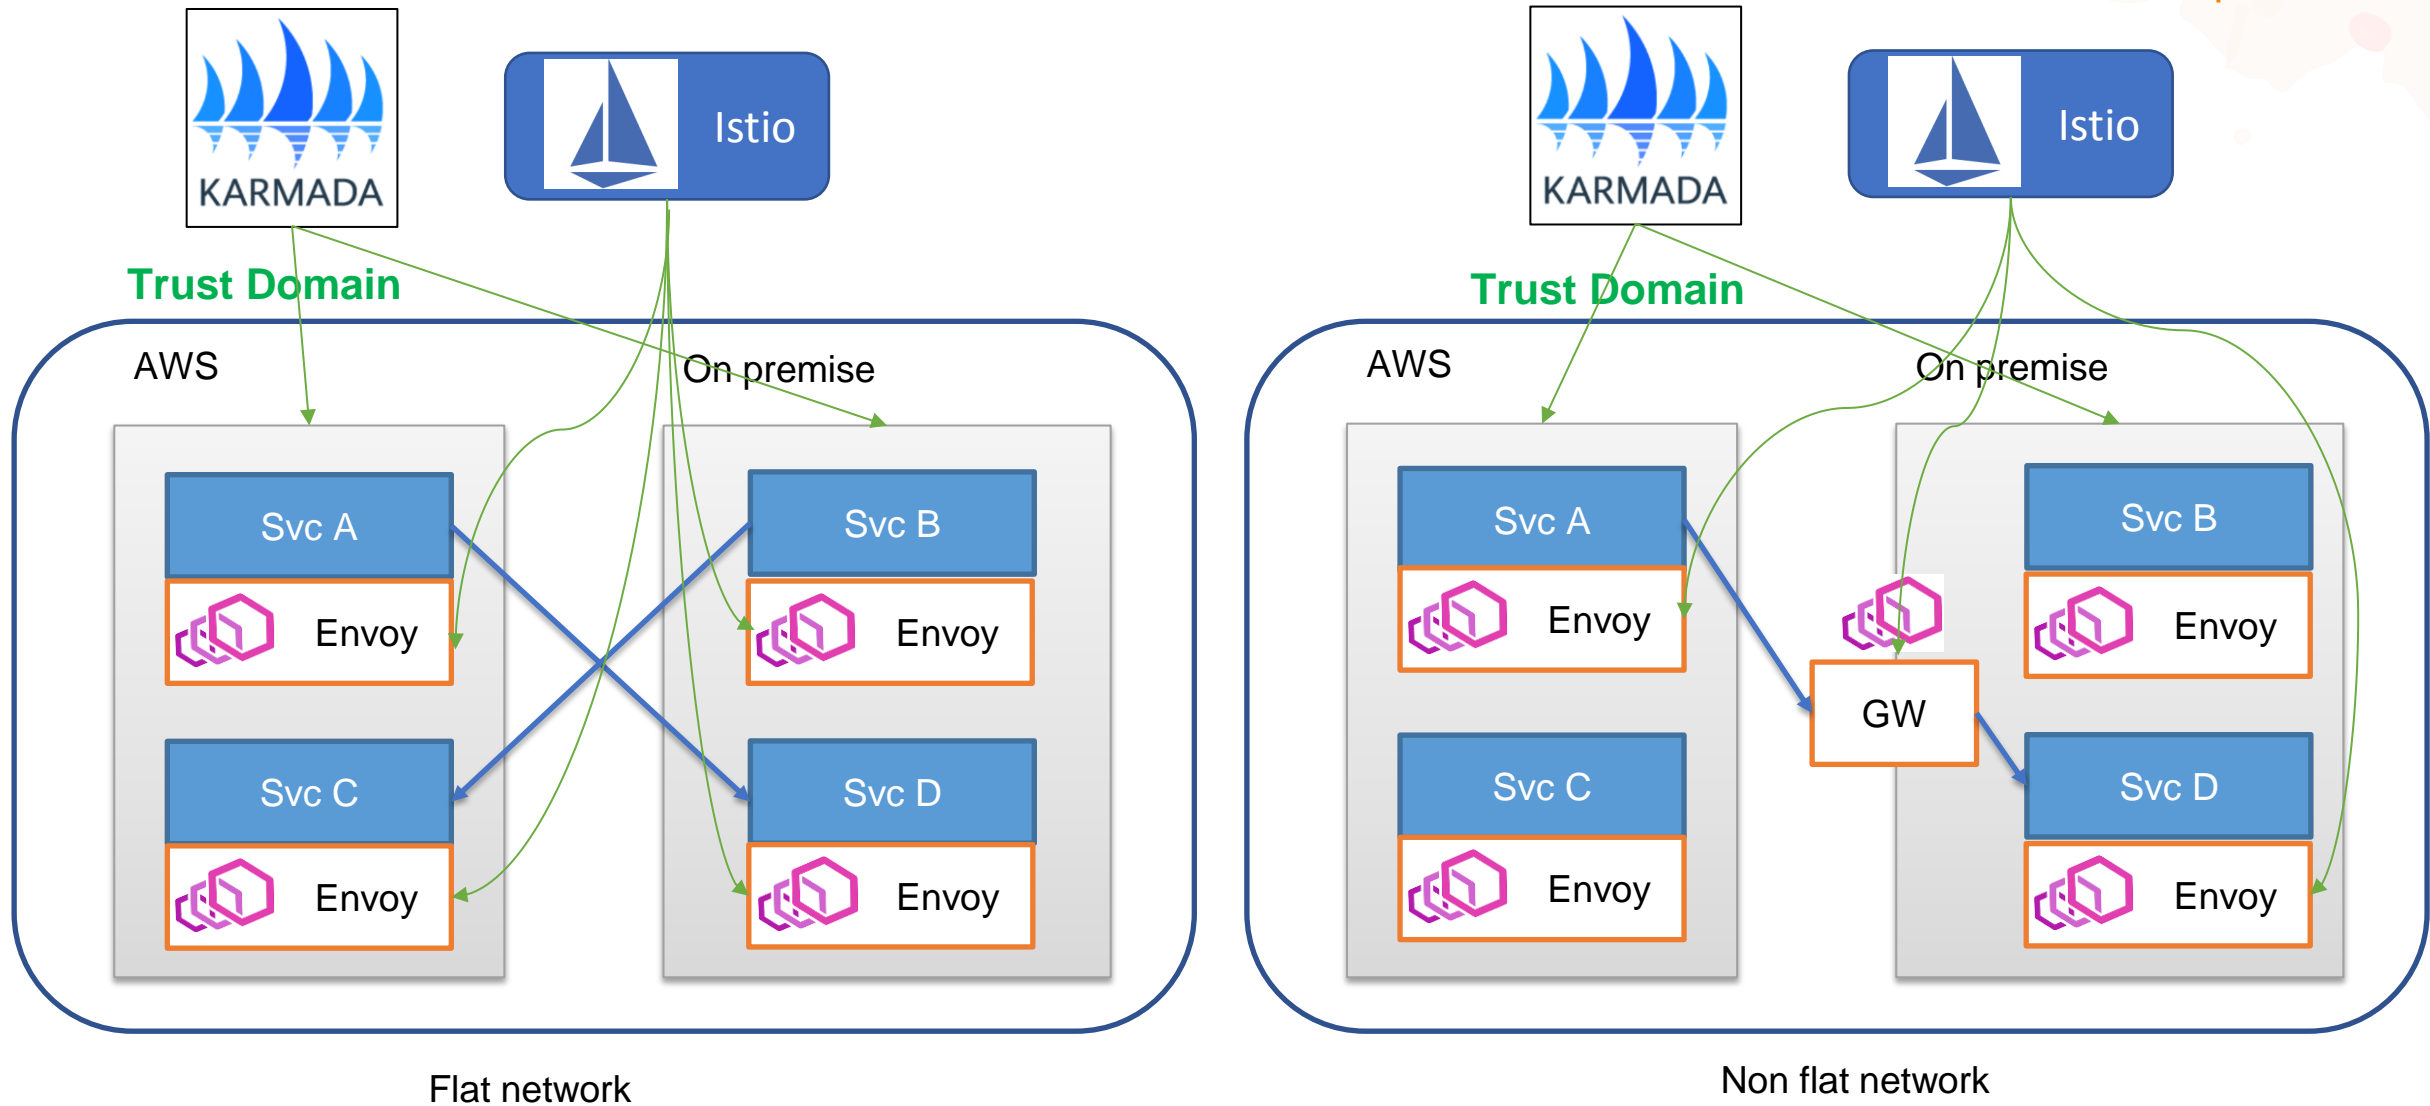




# Multi cloud Traffic Management: Karmada Way

- Submariner to build connect overlay networks of different cloud clusters
- Export and import services between clusters with Multi-cluster Service APIs.
- Integrate with mature service mesh Istio

# Multi cloud Traffic Management: Karmada + Istio



# Multi cloud Traffic Management: DNS Resolution

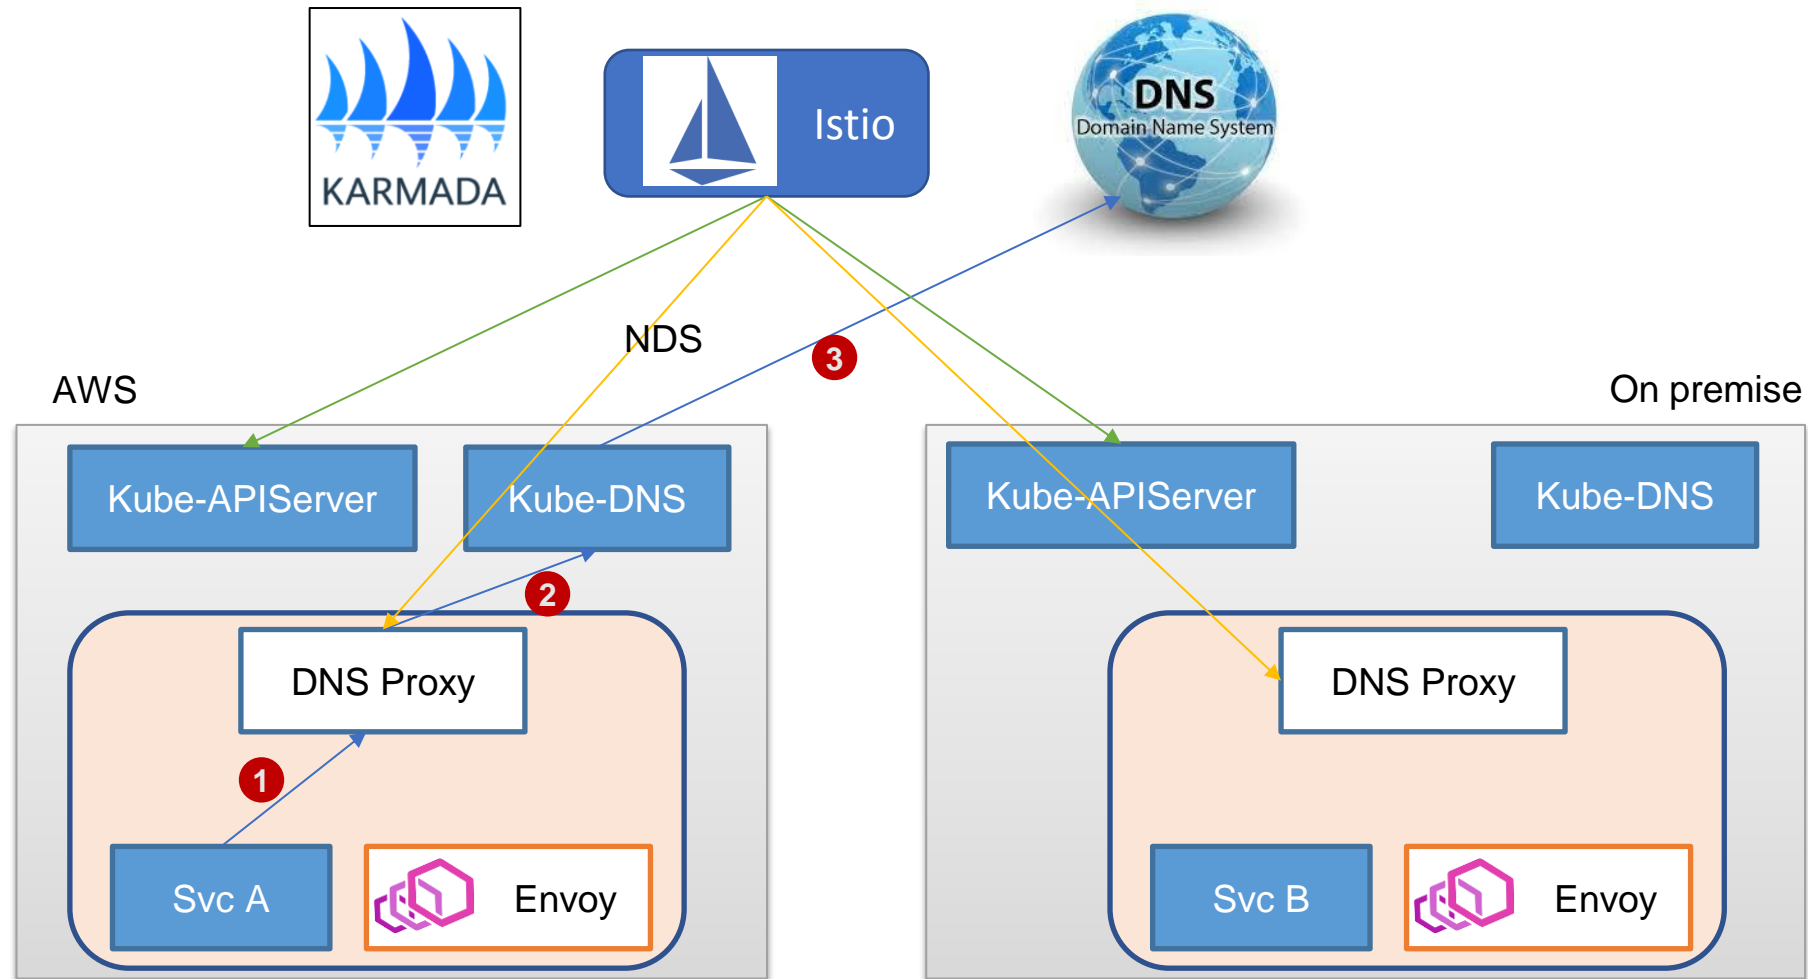


KubeCon



CloudNativeCon

Europe 2022



# Multi cloud Traffic Management: Flat network



KubeCon



CloudNativeCon

Europe 2022

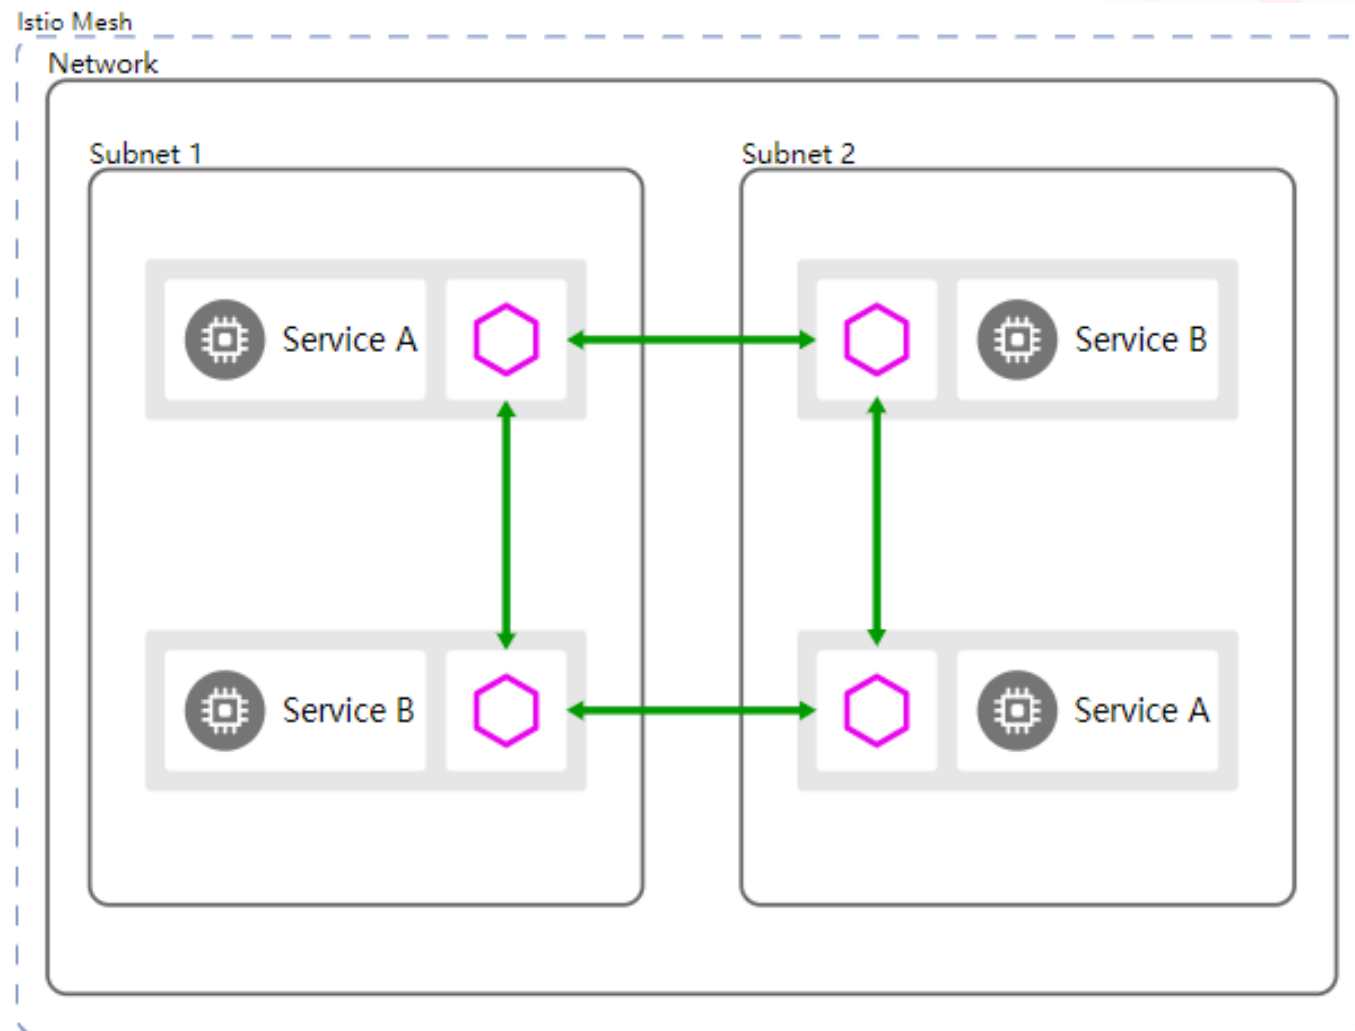
## Flat Network

### Pros

- ① Low latency east-west traffic, as gateway is not needed
- ② Inherit all capabilities from single cluster

### Cons

- ① Complexity: need additional tool to build flat network
- ② Security: not secure as all workloads are within a single network.
- ③ No overlapping service ip ranges



# Multi cloud Traffic Management: Different networks



KubeCon



CloudNativeCon

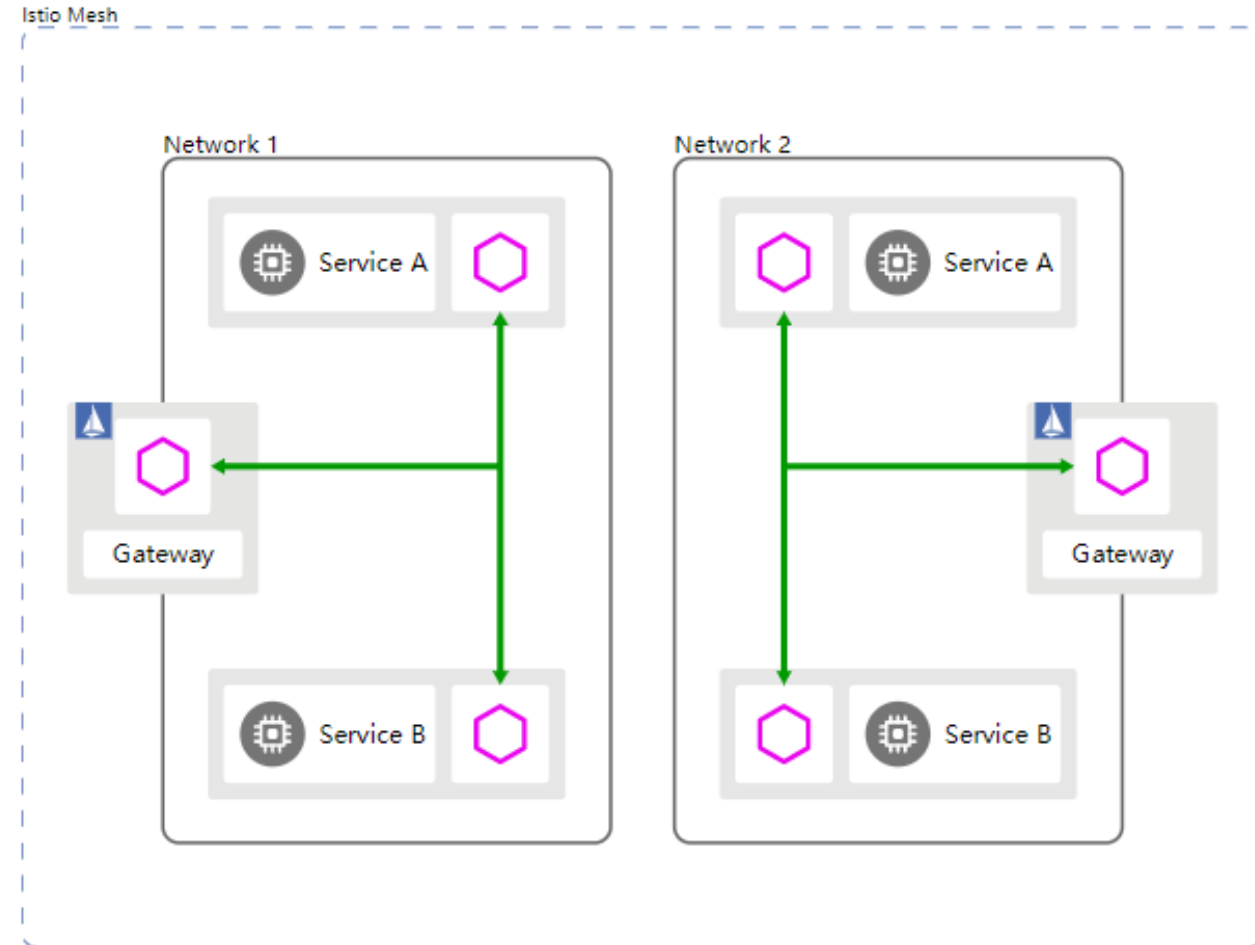
Europe 2022

## Pros

- ① Isolation: network segmentation
- ② Security: east-west traffic is encrypted and through gateway
- ③ Scaling of network addresses
- ④ Cost saving

## Cons

- ① Cross network service communication requires east-west gateway
- ② Gateway works in TLS AUTO\_PASSTHROUGH mode, lack of routing capabilities as flat network
- ③ Additional hops.



# Conclusion

- Multi cloud evolution and what's challenged it bring about
- What Karmada can do for multi cloud
- Inter-Cloud communication with Istio



KubeCon



CloudNativeCon

Europe 2022



KubeCon



CloudNativeCon

Europe 2022

# QA